
Winning Tickets Are Not Posterior Modes: Evidence from Posterior Support and Magnitude Controls

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The lottery ticket hypothesis suggests that dense neural networks contain sparse
2 subnetworks that can be trained in isolation when rewound to their original ini-
3 tialization. A tempting Bayesian interpretation is that these tickets correspond to
4 posterior modes, or to posterior-basin supports. We turn this interpretation into
5 a pre-specified, falsifiable test: posterior-induced masks must match IMP tickets
6 better than local magnitude, dense-trajectory, and pruning-process controls, not
7 merely better than uniform random masks. Across five-seed MNIST and Fashion-
8 MNIST sparsity sweeps and a CIFAR-10 ResNet-20 epoch-1 rewind setting, four-
9 teen posterior approximations—SGLD, SGHMC, cyclical and parallel-tempered
10 SGLD, SWAG, the Laplace family up to joint-group full covariance, subspace
11 HMC, and deep ensembles—fail this stronger criterion. Posterior supports often
12 beat random masks, but they remain tied to chain-start magnitude, dominated by
13 dense magnitude, or closer to rewind controls than to IMP. Support-equivalence is
14 graded rather than binary: only a rank-128 Hessian-plus-diagonal Laplace crosses
15 the mask-overlap threshold, and even it fails the layer-sparsity and basin-count
16 gates, so a posterior account survives on one axis at one fidelity. We then give a
17 positive account of what winning tickets are: trajectory and pruning-process ob-
18 jects. Dense-trajectory magnitude explains most of the fixed-mask training gap
19 and a process-selected IMP-only residual explains the rest, with an extensive bat-
20 tery of process controls failing to substitute for it. This account is consistent with
21 axial-subspace views of IMP and inconsistent with a posterior-mode view. Ex-
22 act dense full-network full-covariance CIFAR posterior baselines remain an open
23 limitation.

1 Introduction

24 The lottery ticket hypothesis (LTH) states that a dense randomly initialized network contains sparse
25 subnetworks that can be trained to match the dense network when reset to their original initialization
26 [5]. This result raises a basic question about the object discovered by iterative magnitude pruning
27 (IMP): is a winning ticket a combinatorial artifact of the optimization trajectory, or does it reflect a
28 deeper probabilistic structure of the posterior over weights?

30 A Bayesian reading is attractive: SGD and its noisy variants connect to approximate Bayesian in-
31 ference [19, 13], and Bayesian deep learning emphasizes that useful solutions are distributed across
32 broad posterior regions [20, 8]. If winning tickets correspond to posterior modes or posterior-basin
33 supports, LTH would be evidence for a selection view of learning: pruning would reveal posterior
34 structure already present in the trained model.

35 This paper tests that interpretation directly. We ask:

36 Do posterior samples or posterior basins induce sparse supports that explain IMP
37 winning-ticket masks beyond dense magnitude, initialization magnitude, rewind
38 magnitude, chain-start magnitude, and pruning-process controls?

39 The paper makes three first-class contributions. **(i) A controlled negative empirical result.** We state
40 the posterior-mode explanation as a *pre-specified*, falsifiable support-equivalence claim—posterior-
41 induced supports must beat local magnitude controls, not just random masks, and mode/ticket distri-
42 butions must match in mask and function space—with the operational gates fixed before the poste-
43 rior runs. The gates are then evaluated across repeated MNIST/Fashion-MNIST sweeps and fourteen
44 CIFAR-10 ResNet-20 posterior approximations (stochastic samplers, Laplace variants, SWAG, sub-
45 space HMC, and direct mode/ticket probes), with a locked final-test confirmation that matches three
46 pre-registered numerical predictions. Support-equivalence is *graded* rather than binary: a single
47 rank-128 Hessian-plus-diagonal Laplace passes the mask-overlap axis while still failing the layer-
48 sparsity and basin-count axes. **(ii) A positive trajectory and pruning-process account.** IMP tickets
49 are explained by dense-trajectory magnitude plus a process-selected IMP-only residual rather than
50 posterior-mode supports. Proposition 1 formalizes why high-SNR posterior TopK masks must track
51 the chain-start magnitude support. **(iii) A reusable posterior-vs-IMP audit framework.** The oper-
52 ational support-equivalence gate, the fourteen-approximation harness, the pre-registered locked-test
53 protocol, and the surrounding audit suite are released as a reusable testbed so other posterior fam-
54 ilies can be tested against IMP tickets under the same operational contract. The release adds a
55 standing open-challenge protocol: third parties may submit any posterior construction—including
56 multimodal families we could not exhaust—and we commit to re-running the unmodified gate and
57 registering the outcome.

58 The conclusion is negative but scoped. Posterior masks beat random masks but track the dense-
59 checkpoint or chain-start magnitude support; chains that move do so without moving toward the IMP
60 ticket; high function-space CKA coexists with failed mask-distribution and basin-count criteria; and
61 the strongest process-control signal is the IMP-selected residual. The cost is quantitative: treating
62 IMP tickets as posterior modes underperforms IMP by 2.3 percentage points on the locked full-data
63 CIFAR-10 ResNet-20 row (posterior 0.872 versus IMP 0.895) and by 5.3 percentage points on the
64 non-residual TinyCNN at matched sparsity (posterior 0.738 versus IMP 0.791), a direct accuracy
65 loss for posterior-based pruning pipelines.

66 2 Related Work

67 **Lottery tickets and pruning.** LTH was introduced by Frankle and Carbin [5] and connected to
68 rewinding and linear mode connectivity by Frankle et al. [6]; randomly weighted networks with
69 learned masks emphasized the selection view [16]. Single-shot pruning methods such as SNIP [10]
70 and SynFlow [18] provide non-Bayesian controls.

71 **Mode connectivity and subspaces.** Low-loss connectivity between solutions challenges the view
72 that trained models are isolated modes [7, 2]. For lottery tickets specifically, Paul et al. [15] ar-
73 gue that IMP masks encode axial subspaces—a direct non-Bayesian competitor to a posterior-mode
74 explanation—and permutation symmetries complicate parameter-space mode counting [4, 1], so
75 our experiments include function-space clustering and linear barrier probes in addition to mask
76 overlap. A central aim is to adjudicate the two competitors directly: the posterior-mode and axial-
77 subspace/trajectory readings make opposite predictions about whether posterior supports beat tra-
78 jectory and pruning-process controls, and the gates below separate them.

79 **Bayesian neural networks.** SGLD provides a simple route from SGD to approximate posterior
80 sampling [19]. SGD itself can be interpreted as approximate Bayesian inference in a stationary
81 regime [13]. Practical posterior approximations such as SWAG [12] and exact small-model sampling
82 studies [8] show that posterior geometry is structured and nontrivial. PAC-Bayesian and Bayesian
83 LTH analyses [17, 9] motivate the need for controls: Bayesian language alone does not imply that
84 IMP masks are posterior-mode supports. Notably, Kuhn et al. [9] report that Bayesian-LTH pruning
85 “should rely primarily on magnitude, secondly on standard deviation”—an independent finding
86 already pointing toward the magnitude tracking our gates make operational.

87 **Bayesian and learned-mask pruning.** A separate line of work *optimizes* posterior-like mask dis-
 88 tributions as training objectives: variational dropout [14], hard-concrete L_0 gates [11], and non-
 89 vacuous PAC-Bayes compression bounds [3]. These answer a design question—can a posterior-
 90 shaped objective produce a good sparse network—whereas we test a structural claim: whether the
 91 posterior of a *standard* run already contains the IMP ticket as a mode-level support. Our learned-
 92 mask baselines (Section 5) include variational and hard-concrete L_0 rows precisely so the two ques-
 93 tions are not conflated.

94 Our contribution is therefore not a new posterior approximation or a claim that Bayesian inference
 95 is uninformative for pruning. It is a controlled support-equivalence test: if posterior modes explain
 96 winning tickets, their sparse supports should match IMP tickets after conditioning on dense magni-
 97 tude, rewind magnitude, trajectory magnitude, and pruning-process alternatives.

98 3 Operational Test

99 Let m_{IMP} denote the mask found by IMP at a fixed sparsity. For a posterior sample $\theta^{(s)}$, the simplest
 100 posterior-induced mask is

$$m^{(s)} = \text{TopK}(|\theta^{(s)}|),$$

101 with the same number of kept weights as m_{IMP} . We also construct aggregate posterior masks from
 102 posterior mean, RMS, signal-to-noise ratio, high variance, and low variance.

103 We compare all masks by support Jaccard similarity to m_{IMP} :

$$J(m, m_{\text{IMP}}) = \frac{|m \cap m_{\text{IMP}}|}{|m \cup m_{\text{IMP}}|}.$$

104 The key controls are random same-sparsity, initialization-magnitude, dense trained-magnitude,
 105 rewind-magnitude, and chain-start magnitude masks, SNIP/SynFlow masks, fixed-mask retrain-
 106 ing controls, channel-aligned direct mode/ticket controls, and process-residual interventions. The
 107 posterior-mode interpretation must beat more than random: posterior masks must explain IMP be-
 108 yond the chain-start, dense, rewind, trajectory, and process controls, and posterior samples must
 109 move support in a way that increases IMP alignment.

110 **Strong and weak forms of the hypothesis.** To avoid testing a strawman we separate two readings
 111 of “winning tickets are posterior modes.” The *strong* form (SPMH)—posterior support distribu-
 112 tions coincide with the ticket distribution in mask, layer sparsity, and basin structure jointly—is
 113 already untenable at high SNR by Proposition 1; rejecting it alone would be trivial. The *weak* form
 114 (WPMH), the target of every gate below, asserts only that posterior supports explain IMP tickets
 115 *beyond* the local magnitude, trajectory, and process controls above; it is the operational content
 116 of selection-view readings [16, 20], which require posterior support to carry ticket information the
 117 chain-start magnitude support does not.

118 **Why posterior TopK masks track the chain start.** The chain-start control is not an arbitrary
 119 baseline; a short argument shows it is the control a posterior-sample mask is structurally bound to.
 120 Model the per-coordinate posterior as $\theta_i^{(s)} \sim \mathcal{N}(\mu_i, \sigma_i^2)$, where μ is the posterior mean—in practice
 121 the chain start or MAP estimate—and define the coordinate signal-to-noise ratio $r_i = |\mu_i|/\sigma_i$.

122 **Proposition 1** (High-SNR posterior TopK masks track the chain start). *For any two coordinates*
 123 *with $|\mu_i| > |\mu_j|$, $\Pr[|\theta_i^{(s)}| > |\theta_j^{(s)}|] \rightarrow 1$ as $\min(r_i, r_j) \rightarrow \infty$. Consequently $\text{TopK}(|\theta^{(s)}|) \rightarrow$*
 124 *$\text{TopK}(|\mu|)$, the chain-start magnitude mask, and a posterior-sample mask cannot beat the chain-*
 125 *start control in the high-SNR limit.*

126 *Proof sketch.* Write $|\theta_i^{(s)}| = |\mu_i| |1 + Z_i/r_i|$ with $Z_i \sim \mathcal{N}(0, 1)$. As $r_i \rightarrow \infty$, $Z_i/r_i \rightarrow 0$ in
 127 probability, so $|\theta_i^{(s)}|/|\mu_i| \rightarrow 1$. For coordinates with a fixed magnitude gap the ordering is therefore
 128 preserved with probability tending to one; the only coordinates that can flip across the TopK thresh-
 129 old are those whose mean magnitudes straddle that threshold within $O(\sigma)$, so the expected Jaccard
 130 of $\text{TopK}(|\theta^{(s)}|)$ with $\text{TopK}(|\mu|)$ is governed by the SNR of the boundary coordinates alone. $\square$

A finite-sample form, verified on saved states. Let τ be the K -th largest $|\mu_i|$ and $\rho_i = ||\mu_i| - \tau|/\sigma_i$ the boundary SNR. Each coordinate swapped into $\text{TopK}(|\theta^{(s)}|)$ displaces a kept one, and at the cut τ every swap forces an excluded coordinate above τ or a kept one below it, so $\mathbb{E}[D] \leq B(\tau) = \sum_{\text{kept}} \bar{Q}(\rho_i) + 2 \sum_{\text{excl.}} \bar{Q}(\rho_j)$ ($\bar{Q}$ the normal upper tail). Both masks keep K coordinates, so $J = (K - D)/(K + D)$ is convex in D and Jensen gives $\mathbb{E}[J] \geq (K - B(\tau))/(K + B(\tau))$, recovering Proposition 1 as $\min_i \rho_i \rightarrow \infty$. A generated audit fits the surrogate $(\hat{\mu}, \hat{\sigma})$ to the ten saved SGLD samples per seed: per-seed lower bounds 0.667–0.673 against observed Jaccard means 0.881–0.884, holding in 5/5 seeds (a surrogate-based prediction, not a theorem about SGLD itself).

Proposition 1 predicts the regime we observe: trained weights are high-SNR for the bulk of their large-magnitude coordinates, so the posterior-sample TopK mask coincides with the chain-start mask except at the keep/prune boundary—exactly the measured 0.91–0.98 posterior-to-chain overlap with a posterior-minus-chain gap indistinguishable from zero. It also sharpens the burden of proof: a posterior account can differ from the chain start only at low-SNR boundary coordinates, where its samples are noisiest.

Why rank-bounded posteriors should struggle (informal). Proposition 1 bounds marginal disagreement at the level of single coordinates, but the gate also tests *joint* mask structure. A local Gaussian posterior $q = \mathcal{N}(\theta, UU^\top + D)$ with low-rank factor $U \in \mathbb{R}^{p \times k}$ only induces order- $k + 1$ Gaussian latent dependencies among the boundary coordinates, whereas an R -round deterministic IMP process layers an additional pairwise survival dependence on the boundary set at every round, accumulating high-order dependencies among coordinates that are not Gaussian-latent. The informal expectation is therefore that increasing the covariance rank of a single Gaussian family closes the marginal-tracking axis first and the joint-distribution axes only with much larger rank. We do not formalise this as a theorem—a clean total-variation separation would require a quantitative coupling argument on the TopK boundary set—and treat the locked-budget covariance sweep below as *empirical* evidence consistent with this intuition, not as a verification of a formal separation.

The locked-budget covariance sweep matches this expectation. Of all rank- k Gaussian Laplace rows tested (rank 16/32/64/128 plus joint-group 22,064–270,896-weight exact rows), only the rank-128 row crosses any mask-overlap threshold while the layer-sparsity and basin-count axes still fail, and the exact 270,896-weight joint-group row returns to Hamming overlap 0.000: rank moves the joint mask distribution along only one of three gate axes.

Validation-derived locked-test predictions. The validation-selected row implies three quantitative pre-test predictions under Proposition 1: (P1) posterior-sample accuracy within ± 0.005 of the validation value 0.874; (P2) layer-sparsity KS rejection $p < 0.01$; (P3) posterior-versus-IMP Hamming overlap below 0.05. The single-shot locked test on the held-out test split returned, respectively, 0.872, $p = 0$ to four decimals, and 0.0000—all three derivable predictions confirmed at the same 4/6 threshold-pass count as the validation row. The operational gates themselves (the proposal thresholds in the distribution-probe script) were fixed in code before any CIFAR run; P1–P3 are derived from validation numbers, not externally pre-registered, and are reported as a transparency check rather than a registered protocol.

4 The Posterior-Mode Account Fails the Gates

Mask-overlap metric, made precise. Throughout this section, the “Hamming overlap” threshold reports the empirical overlap of two pairwise-distance histograms—one over posterior mask pairs and one over IMP ticket pairs—rather than the direct posterior-to-ticket nearest-pair distance. Two mask sets whose internal diversity profiles agree can therefore clear the 0.7 threshold even when their cross-set Hamming distance is large. We report the cross-set Hamming mean (“cross_mean”) in every `metrics.json` alongside the histogram overlap to keep the two semantics distinguishable; the headline gate, as the proposal stipulates, is the distribution-equivalence one, and sections interpreting a Hamming-overlap pass additionally rely on the cross-set distance and layer-sparsity KS axes.

4.1 Random Is Too Weak a Baseline

Table 1 shows the current Gate1 sparsity sweeps. Each row aggregates five seeds, three independent dense-start SGLD chains per seed, and posterior masks at the matched IMP sparsity.

Table 1: MNIST and Fashion-MNIST 5-seed Gate1 sparsity sweeps. Posterior masks beat random but fail against stronger controls at every sparsity. “Post-Chain” is the Jaccard overlap between posterior masks and their chain-start magnitude masks. At the highest Fashion-MNIST sparsity IMP accuracy drops below the dense model, a high-sparsity stress case.

Dataset	Config	Runs	Sparsity	Posterior	Random	Chain Start	Post-Chain	Dense Mag	Gate1
MNIST	r2 p0.30	5	0.5100	0.3539	0.3246	0.3537	0.9303	0.8259	Fail
MNIST	r3 p0.30	5	0.6570	0.2669	0.2071	0.2670	0.9115	0.7258	Fail
MNIST	r5 p0.30	5	0.8319	0.2373	0.0918	0.2385	0.9298	0.5924	Fail
MNIST	r8 p0.30	5	0.9424	0.1275	0.0298	0.1275	0.9423	0.4028	Fail
Fashion	r2 p0.30	5	0.5100	0.3516	0.3246	0.3516	0.9314	0.7385	Fail
Fashion	r3 p0.30	5	0.6570	0.2595	0.2070	0.2594	0.9160	0.6378	Fail
Fashion	r5 p0.30	5	0.8319	0.2114	0.0917	0.2122	0.9212	0.5255	Fail
Fashion	r8 p0.30	5	0.9424	0.1300	0.0297	0.1300	0.9408	0.3813	Fail




Figure 1: Gate1 control comparison on MNIST and Fashion-MNIST. Posterior masks beat random masks, but they track chain-start magnitude and remain far below dense magnitude across sparsities.

182 Posterior masks beat random support at every tested MNIST/Fashion-MNIST sparsity, but the ran-
183 dom gap is explained by chain-start magnitude: posterior-minus-chain-start gaps stay near zero at
184 0.91–0.94 posterior-to-chain-start overlap, and “near zero” is equivalence-tested (a TOST reanalysis
185 bounds the gap within ± 0.005 Jaccard on all eight Gate1 rows). Dense magnitude is much closer to
186 IMP (dense-minus-posterior gaps 0.275–0.472 on MNIST, 0.251–0.387 on Fashion-MNIST). The
187 generated mode-distribution audit shows the same pattern: posterior supports beat uniform random
188 masks in 58 of 59 grouped comparisons, while none of the 59 posterior-versus-chain-start compar-
189 isons beats the matched control by more than 0.005 Jaccard (43 are practically tied to the control,
190 15 favor chain-start magnitude, and one is mixed) and rewind magnitude is closer to the IMP ticket
191 than posterior support in 55 of 57 grouped comparisons.

192 4.2 Sampler Movement Does Not Point Toward Tickets

193 The CIFAR-10 ResNet-20 evidence tests whether the MNIST/Fashion failure is an artifact of small
194 models or weak image baselines. In the long-budget epoch-1 rewind row, dense accuracy is about
195 88.5% and IMP reaches about 89.8% at 83.2% sparsity, yet SGLD, SGHMC, cyclical SGLD,
196 SWAG, diagonal Laplace, KFAC-style Laplace, low-rank Hessian-plus-diagonal Laplace, and sub-
197 space HMC all fail the same control. Increasing sampler scale often lowers posterior-to-chain-start
198 overlap, but posterior-to-IMP overlap does not improve and the epoch-1 rewind support remains
199 closer to IMP.

200 The negative movement result is broad: the sampler grid spans dense-start and independent-start
201 multi-chain SGLD and SGHMC, cyclical SGLD, a 20-snapshot SWAG row, diagonal and KFAC-
202 style Laplace, rank-16/32/64/128 low-rank Laplace, exact head and selected/tensor-block/joint-
203 group full-covariance Laplace, and low-dimensional HMC in random, trajectory, and top-Hessian
204 subspaces. Full-batch HMC and an exact dense full-network Laplace check on small digits MLPs
205 validate the higher-fidelity code paths while preserving the same support failure.

206 4.3 Richer Covariance Still Fails the Support Gate

207 The covariance objection is the strongest bounded limitation, so we push exact covariance as far
208 as the workstation budget permits: five exact block-diagonal and joint-group rows covering 22,064,
209 68,144, 68,144, 86,576, and 270,896 trainable weights. Every row tells the same story: sample
210 accuracy stays near 88%, global posterior-to-chain-start support falls as fidelity grows (from 0.8287




Figure 2: CIFAR-10 ResNet-20 epoch-1 rewind movement diagnostics. Across SGLD, SGHMC, cyclical SGLD, full-network SWAG, diagonal Laplace, KFAC-style Laplace, and low-rank Hessian-plus-diagonal Laplace, moving support away from the chain start does not increase overlap with the IMP mask.

at the 22,064-parameter row down to 0.7148–0.7863) so the posterior does move, yet the selected-block posterior-minus-chain gain stays at or below zero (between -0.0114 and $+0.0036$) and the rewind control remains closer to IMP by 0.029–0.036 throughout. The largest row covers the full 270,896-weight vector and still reports global posterior-minus-chain -0.0019 . These rows remove the weight-tensor coverage objection, while literal dense CIFAR full-network covariance remains outside budget (a float64 dense matrix needs 553.1 GiB; with Cholesky, 1,106.3 GiB). Per-row numbers are in the ledger.

4.4 Function Space Agreement Is Not Mask Equivalence

We also make the original proposal’s distributional criterion explicit. On digits, posterior sample masks fail two of the four stated thresholds: layer-sparsity KS has $p = 0.0788$ and pairwise mask-Hamming overlap is 0.631, while logit CKA passes. Mean-shift clustering in raw parameter PCA collapses posterior samples to one representative mode; the raw-parameter basin entropy is 0.0 nats with effective cluster count 1.0. The full-data CIFAR-10 ResNet-20 strengthens this separation: posterior samples can keep logit/final-hidden activation CKA above 0.93/0.91 while failing layer-sparsity, Hamming, or posterior-basin-count criteria.

Channel alignment does not change this conclusion: activation-channel and weight-correlation alignment both still fail layer KS and Hamming while retaining high activation CKA; a block-coordinate ResNet channel audit leaves global-channel posterior/ticket Hamming near 0.21 rather than ticket-like agreement; and an exact stage-1 enumeration validates the local solver on a small ResNet subgraph while exhaustive full-data graph isomorphism remains infeasible. A permutation-invariant metric closes this objection without enumerating permutations: per-tensor entropic Gromov–Wasserstein distances over channel graphs are relabeling-invariant by construction, and independent per-tensor couplings are a relaxation strictly more favorable to the posterior than any valid network symmetry. On the saved full-data SGLD artifact the metric is validated (a genuine producer/consumer channel permutation stays at the self-distance floor despite large coordinate Hamming distance), yet posterior samples stay at cross-seed ticket-diversity distance from their own-seed ticket in 5/5 seeds ($0.96\text{--}1.06\times$ the cross-ticket mean): relabeling moves posterior masks no closer to the ticket at all.

The multi-chain and high-fidelity direct rows address the concern that a single local chain was too weak: a dense-start direct probe with 75 cyclical-SGLD posterior samples, and an independent-start version with 15 dense chain starts, both fail layer KS and Hamming with high CKA. A streamed exact joint-group direct probe over all 270,896 weight parameters closes the full-weight Gaussian gap: 25 joint-group posterior samples keep sample accuracy at 88.3%, posterior-to-chain-start Hamming is 0.0503, layer-sparsity KS is $p = 1.1 \times 10^{-8}$, Hamming overlap is 0.000, logit/final-hidden activation CKA remains 0.937/0.920, and all 25 joint-group samples still collapse to one parameter-PCA basin.

Multimodal approximations: deep ensembles and tempered chains. Every approximation above is a local or single-chain family, so we add the two standard multimodal constructions. A *deep ensemble* of ten independently initialized full trainings—the approximation Wilson and Izmailov [20] advocate—keeps accuracy (0.880 versus dense 0.882) yet fails layer-sparsity KS ($p = 5.3 \times 10^{-7}$) and Hamming overlap (0.000) while passing logit CKA (0.931), members below IMP 5/5; *parallel-tempered SGLD* (1–8 ladder, 4–18% swap acceptance) fails identically (KS $p = 4.7 \times 10^{-9}$, overlap 0.000, CKA 0.935). Multimodality alone does not move posterior supports toward IMP tickets.

Exploratory robustness: non-residual architecture (seed-level cell). A reviewer could attribute the failure to residual architecture. We treat this cell as an *exploratory robustness cell* rather than a confirmatory row: it evaluates on the test split directly and bounds the architecture axis only; the dataset axis is handled by the validation-selected CIFAR-100 locked test below. We run the seed-level part of the matched probe on a non-residual three-convolution TinyCNN (no skip connections, no global average pooling; we retain BatchNorm so this is a non-residual—but not BN-free—robustness point) at the same five-round $p=0.30$ schedule and 0.832 sparsity, reporting the per-seed direction checks and posterior-to-chain-start Hamming (the activation-CKA stage stalls on TinyCNN’s wide 4096-dimensional pre-classifier feature, so the joint-distribution gate axes were not computed). The seed-level direction is unchanged: IMP beats the dense chain start in 5/5 seeds (0.791 versus 0.781); SGLD posterior samples stay below both in 5/5 seeds (0.738); and posterior TopK masks track the chain start with posterior-to-chain-start Hamming 0.061 (95% CI [0.060, 0.062]); the joint-distribution axes for a non-residual architecture remain an open robustness item.

Dataset robustness: CIFAR-100 with a validation-selected locked test. The dataset axis is tested at two levels on CIFAR-100 ResNet-20-w16 at the same $p=0.30$ schedule and 0.832 sparsity. First, an exploratory five-seed direct probe on the full training set (test-split evaluation): the seed-level direction holds (IMP beats dense 5/5 at 0.636 versus 0.629, posterior samples below both 5/5 at 0.617); on this harder task IMP supports coincide with dense magnitude (chain-versus-tickets Hamming overlap 1.000), so posterior masks clear the Hamming-overlap axis (0.816) while still failing layer-sparsity KS ($p = 1.3 \times 10^{-6}$) and activation CKA—4/6, as the rank-bounded intuition (§3) predicts. Second, the CIFAR-10 validation/locked-test protocol: a held-out 10% validation split selected the row and a single locked final-test rerun reproduced it exactly (layer KS $p = 9.0 \times 10^{-8}$, Hamming overlap 0.000, logit CKA 0.884, 3/6; IMP beats dense 5/5 at 0.620 versus 0.596, posterior 5/5 below both at 0.584—a 3.6-point gap to IMP). Under proper split separation IMP no longer collapses onto dense magnitude and the Hamming axis fails again, so the locked CIFAR-100 result is *stronger* than the exploratory probe suggested.

Support-equivalence is graded, not binary. The criterion behaves as a graded property in covariance fidelity rather than a single pass/fail switch. Of every posterior approximation tested, exactly one—the rank-128 Hessian-plus-diagonal Laplace direct probe—crosses any gate axis: its 50 posterior samples pass the Hamming-overlap threshold (0.816) with high logit/final-hidden activation CKA (0.932/0.910), but still fail the layer-sparsity KS threshold ($p = 2.0 \times 10^{-6}$) and collapse to one parameter-PCA basin. The pass is replicable, not single-run noise: an independent five-seed group (seeds 5–9) under the identical protocol reproduces the full profile (Hamming overlap 0.749, layer KS $p = 1.5 \times 10^{-6}$, one basin, same 5/6 axes). Neighbouring fidelities bracket it: rank-32/64 movement rows do not cross the mask-overlap threshold and the exact 270,896-weight joint-group row returns to Hamming overlap 0.000. A posterior account therefore has reproducible traction in exactly one narrow band—one family, one rank, one of three gate axes—reported as a partial-equivalence regime, not as support for the hypothesis: a replicated single-axis pass is still not the distributional mask/basin equivalence the hypothesis requires.

The layer-sparsity KS values in these direct rows are descriptive diagnostics over correlated posterior samples, not independent sample-level significance tests. The saved full-data SGLD artifact permits a seed-level paired reconstruction: posterior samples are farther from the same-seed IMP ticket than the chain-start support in 5/5 raw seeds and 5/5 activation-aligned seeds, and an independent seeds 5–9 extension run with saved masks reproduces the direction in 5/5 further seeds—10/10 across both groups, two-sided sign $p = 0.002$. The cyclical-SGLD multichain, rank-128 low-rank Laplace, and




Figure 3: CIFAR-10 ResNet-20 dense-trajectory controls. Magnitude supports along the matched dense trajectory, especially around epochs 10–30, and the best aggregate trajectory magnitude score are closer to the IMP mask than any global posterior support in the movement diagnostics.

joint-group Laplace direct rows also write per-seed `mask_artifacts.npz` for the same on-demand paired audit; pooled-sample KS p-values for them remain descriptive diagnostics.

5 What Winning Tickets Are: A Trajectory and Process Account

The posterior-mode account fails the gates, but a winning ticket still demands a positive explanation. We give one: a winning ticket is a *trajectory and pruning-process object*, with a trajectory-magnitude component that any dense run exposes and a process-selected IMP-only residual that only iterative pruning produces. Neither part is recovered by posterior structure.

Dense trajectory magnitude is the first part, and it is useful: in the matched CIFAR dense-trajectory probe, trajectory magnitude supports are closer to IMP than posterior supports, and fixed-mask training shows that final dense, epoch-10, and RMS trajectory masks are much better than random. They are still below IMP, so the remaining question is what IMP adds.

Residual swap and removal-order controls localize the gap to a process-selected IMP-only residual. Anatomy, predictor-mask, cross-seed, direct-transfer, base-compatibility, posterior-ordering, and stratified controls give the same localization. Learned mask baselines do not replace that residual: Gem-Miner-style, Bernoulli/Concrete variational pruning, and the hard-concrete L0 gate baseline are below IMP in the current CIFAR support/accuracy rows and do not rescue the calibration/OOD tradeoff. Their learned objective can change uncertainty behavior, but not into a ticket-support explanation.

The strongest process controls test whether simpler score composition explains the residual. A tensor-matched excluded oracle stays below the round-selected IMP residual, and a tensor+score-matched oracle remains below it even after matching parameter tensor and within-tensor round-score deciles. Three residualization checks then confirm the residual is not a posterior artifact. A residualized round-score projection that removes base, dense-final, and final-IMP magnitude subspaces reduces both accuracy and oracle overlap; removing the diagonal-Laplace posterior score subspace (posterior RMS, standard deviation, and RMS-minus-dense features) leaves a 0.23-point round-vs-residualized accuracy gap (5/5 positive seeds, 95% CI $[-0.03, 0.48]$) with oracle overlap dropping 0.1923; and removing a learned trajectory/process subspace leaves a 0.48-point gap (5/5 seeds, 95% CI $[0.29, 0.67]$) with overlap dropping 0.1890. The learned process subspace captures broad directions but does not substitute for the particular residual coordinates IMP selects.

Positive prediction: cross-seed mask transfer. The trajectory and pruning-process account makes two falsifiable predictions that a posterior-mode account (which predicts seed-invariant support, hence little within/cross gap) does not: (P+1) same-seed IMP masks beat masks transferred from another seed, because the trajectory selects support per seed; (P+2) cross-seed masks still carry above-random support precision, because the trajectory-magnitude component is partially shared. Both hold in the five-seed transfer probe: same-seed accuracy minus IMP is -0.0090 ($[-0.0111, -0.0068]$) versus cross-seed -0.0235 ($[-0.0269, -0.0201]$), a 1.5-point drop (P+1);

and cross-seed added support precision is 0.2388 ([0.2334, 0.2442]), roughly twice the random baseline 0.1264 (P+2).

6 Discussion

These results constrain a simple Bayesian interpretation of lottery tickets. The posterior signal exists in a weak sense—posterior-induced masks are not random—but the same signal is already present in the dense, rewind, trajectory, or chain-start magnitude support. The explicit distribution-equivalence audit turns this into a negative result rather than a single-metric impression, and the evidence favors a trajectory/magnitude subspace account plus a process-selected residual over a posterior-mode account.

This adjudicates the two competitors of Section 2 in favor of the axial-subspace reading. Paul et al. [15] hold that IMP masks encode trajectory axial subspaces, not posterior structure; the posterior-mode reading predicts the opposite. Our gates test that prediction and it fails: trajectory magnitude and the IMP residual carry the ticket signal, posterior supports do not, and the graded rank-128 exception passes only one axis at one fidelity. The evidence is consistent with the axial-subspace account and inconsistent with a posterior-mode account in the tested settings.

This conclusion is compatible with mode-connectivity results, and a generated linear connectivity audit makes the point concrete: MNIST and Fashion-MNIST dense-to-IMP barriers are nearly zero (0.0026 and 0.0395) while posterior support stays below the chain-start control, and CIFAR-10 long SGLD and SWAG barriers are large (3.0827 and 3.7402) while posterior support stays tied to chain-start support. Linear barriers are orthogonal landscape diagnostics, not evidence of posterior-ticket equivalence.

Scope of the claim. The claim is not that Bayesian neural-network posteriors are useless, nor that no posterior approximation could ever correlate with IMP coordinates. The confirmatory evidence in this paper is the five-seed MNIST/Fashion-MNIST sparsity sweeps, the CIFAR-10 ResNet-20 epoch-1-rewind direct probe across fourteen posterior approximation configurations, and the validation-selected CIFAR-100 ResNet-20-w16 locked-test row; the non-residual TinyCNN seed-level cell is an additional robustness check under the same matched-sparsity schedule. The claim is narrower: under the posterior families and CIFAR-10/MNIST/Fashion-MNIST settings tested here, posterior-induced sparse supports do not satisfy the support-equivalence claim once dense, rewind, trajectory, and pruning-process controls are included; useful ticket structure is better explained by trajectory magnitude plus a process-selected IMP residual.

7 Limitations and Next Experiments

The evidence is broad, but the paper should still be read as a controlled negative result under the posterior approximations we can currently test, not as a theorem about all Bayesian neural-network posteriors. We also note an explicit multiple-comparisons exposure: confirmatory rows pool five seeds against four to six thresholds across fourteen posterior configurations, so individual pooled p-values are descriptive diagnostics. A generated family-wise audit quantifies the exposure over the sixteen distinct full-data direct configurations ($16 \times 6 = 96$ trials): all sixteen layer-sparsity KS rejections survive Holm–Bonferroni at $\alpha = 0.05$; the two single-axis Hamming-overlap passes (rank-128 and CIFAR-100) match the expected single-axis pass count under pooled per-axis base rates; and no configuration passes the full joint gate. A TOST reanalysis at the paper’s own materiality margin ($\delta = 0.005$ Jaccard) equivalence-passes all eight Gate1 tracking rows, with all 35 CIFAR movement rows one-sided non-superior at the same margin; the four independent seed-level direction checks (saved-artifact SGLD seeds 0–4 and 5–9, TinyCNN, CIFAR-100) jointly have global-null probability 9.5×10^{-7} , and the headline SGLD row alone spans ten seeds (10/10, two-sided sign $p = 0.002$). The conclusion thus rests on family-wise-corrected rejections, equivalence-tested tracking, and joint seed-level direction, not on any single pooled p-value. The main remaining limitations are:

- alignment robustness. Activation-channel and weight-correlation alignment, record-level post-hoc matching, the block-coordinate global channel audit, and exact stage-1 enumeration are negative or bounded; exhaustive full-data graph isomorphism remains infeasible and unimplemented;

- validation/test separation. The selected CIFAR-10 SGLD direct row was rerun on a held-out 10% validation split and a locked final-test rerun has been completed (posterior-sample accuracy 0.872, 4/6 thresholds, matching the validation row); the CIFAR-100 row follows the same protocol (validation-selected, locked test reproducing it at 3/6);
 - BatchNorm posterior policy. ResNet posterior samplers include BatchNorm buffers in sampled state dictionaries under the default train-buffer policy; the full CIFAR SGLD/cSGLD sweep over freeze, recalibrate, and dense-buffer policies keeps the negative direction in 5/6 rows, with frozen-BN vanilla SGLD reported as an invalid sampler setting;
 - direct distribution statistical unit. Pooled direct mode/ticket p-values over posterior samples are descriptive over correlated samples; the SGLD, cSGLD multichain, rank-128 Laplace, and joint-group Laplace direct rows all write per-seed `mask_artifacts.npz` for seed-level paired audits, and the headline SGLD direction now spans ten seeds (sign $p = 0.002$);
 - full-network exact/full-covariance Laplace or another still broader high-fidelity posterior baseline: the covariance ladder up to all-weight joint groups, 20-snapshot SWAG, and subspace HMC is covered, but the exact dense full-network CIFAR posterior remains absent;
 - broader learned-mask distribution variants. Gem-Miner-style, variational, and hard-concrete L0 gates do not rescue support, accuracy, calibration, or CIFAR-100 OOD behavior; additional permutation-aware learned-mask distributions would make this negative claim cleaner;
 - additional process-causality controls. The residual-swap, anatomy, predictor-mask, transfer, ordering, stratification, IMP-process, tensor+score-matched round-exclusion, residualized-score projection, posterior-residualized, and learned-subspace residualized controls do not exhaust the possible causal interventions on the pruning process;
 - Reproducibility Statement. The repository ships an environment lock, Makefile gates, artifact verifier, CI workflow, release archive, source snapshot, verification containers, and an external-validation runbook; public archive, public repository, and external CI receipts are recorded against the current archive SHA and source commit, and the one remaining external receipt is an independent GPU-container run.
- Since the aligned, multi-chain, Laplace-family, locked final-test, and BatchNorm-policy CIFAR runs preserve the current pattern, the next decisive empirical gap is an even stronger high-fidelity posterior baseline; if it stays near the chain-start and rewind controls, the negative result is the right paper, otherwise the claim should weaken to a partial-equivalence statement (the rank-128 row passes 5/6 thresholds and is the strongest example in the current evidence).

421 Ethics Statement

422 This work uses standard public benchmark datasets and does not introduce human subjects data,
 423 private personal data, surveillance data, or safety-critical deployment claims. The main ethical risk
 424 is scientific overclaiming: reading the negative result as a theorem rather than a controlled empirical
 425 result under the tested architectures, approximations, sparsities, and budget. We mitigate it by stating
 426 the operational criterion, reporting negative and bounded findings, preserving open limitations, and
 427 packaging code and artifacts for reproduction. Final submission still requires the human authors to
 428 review and acknowledge the applicable ICLR Code of Ethics.

References

- [1] Samuel K. Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. Git re-basin: Merging models modulo permutation symmetries. In *International Conference on Learning Representations*, 2023.
- [2] Felix Draxler, Kambis Veschgini, Manfred Salmhofer, and Fred A. Hamprecht. Essentially no barriers in neural network energy landscape. In *International Conference on Machine Learning*, 2018.
- [3] Gintare Karolina Dziugaite and Daniel M. Roy. Computing nonvacuous generalization bounds for deep (stochastic) neural networks with many more parameters than training data. In *Conference on Uncertainty in Artificial Intelligence*, 2017.
- [4] Rahim Entezari, Hanie Sedghi, Olga Saukh, and Behnam Neyshabur. The role of permutation invariance in linear mode connectivity of neural networks. In *International Conference on Learning Representations*, 2022.
- [5] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Finding sparse, trainable neural networks. In *International Conference on Learning Representations*, 2019.
- [6] Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, and Michael Carbin. Linear mode connectivity and the lottery ticket hypothesis. In *International Conference on Machine Learning*, 2020.
- [7] Timur Garipov, Pavel Izmailov, Dmitrii Podoprikin, Dmitry P. Vetrov, and Andrew Gordon Wilson. Loss surfaces, mode connectivity, and fast ensembling of dnns. In *Advances in Neural Information Processing Systems*, 2018.
- [8] Pavel Izmailov, Sharad Vikram, Matthew D. Hoffman, and Andrew Gordon Wilson. What are bayesian neural network posteriors really like? In *International Conference on Machine Learning*, 2021.
- [9] Nicholas Kuhn, Arvid Weyrauch, Lars Heyen, Achim Streit, Markus Götz, and Charlotte Debus. Bayesian lottery ticket hypothesis, 2026.
- [10] Namhoon Lee, Thalaiyasingam Ajanthan, and Philip H. S. Torr. SNIP: Single-shot network pruning based on connection sensitivity. In *International Conference on Learning Representations*, 2019.
- [11] Christos Louizos, Max Welling, and Diederik P. Kingma. Learning sparse neural networks through l_0 regularization. In *International Conference on Learning Representations*, 2018.
- [12] Wesley J. Maddox, Pavel Izmailov, Timur Garipov, Dmitry P. Vetrov, and Andrew Gordon Wilson. A simple baseline for bayesian uncertainty in deep learning. In *Advances in Neural Information Processing Systems*, 2019.
- [13] Stephan Mandt, Matthew D. Hoffman, and David M. Blei. Stochastic gradient descent as approximate bayesian inference. *Journal of Machine Learning Research*, 2017.
- [14] Dmitry Molchanov, Arsenii Ashukha, and Dmitry Vetrov. Variational dropout sparsifies deep neural networks. In *International Conference on Machine Learning*, 2017.
- [15] Mansheej Paul, Feng Chen, Brett W. Larsen, Jonathan Frankle, Surya Ganguli, and Gintare Karolina Dziugaite. Unmasking the lottery ticket hypothesis: What’s encoded in a winning ticket’s mask? In *International Conference on Learning Representations*, 2023.
- [16] Vivek Ramanujan, Mitchell Wortsman, Aniruddha Kembhavi, Ali Farhadi, and Mohammad Rastegari. What’s hidden in a randomly weighted neural network? In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020.
- [17] Keitaro Sakamoto and Issei Sato. Analyzing lottery ticket hypothesis from pac-bayesian theory perspective. In *Advances in Neural Information Processing Systems*, 2022.
- [18] Hidenori Tanaka, Daniel Kunin, Daniel L. K. Yamins, and Surya Ganguli. Pruning neural networks without any data by iteratively conserving synaptic flow. In *Advances in Neural Information Processing Systems*, 2020.
- [19] Max Welling and Yee Whye Teh. Bayesian learning via stochastic gradient langevin dynamics. In *International Conference on Machine Learning*, 2011.
- [20] Andrew Gordon Wilson and Pavel Izmailov. Bayesian deep learning and a probabilistic perspective of generalization. In *Advances in Neural Information Processing Systems*, 2020.

477 A Anticipated Reviewer Objections and Responses

478 This appendix collects the five reviewer objections we judge most likely for a negative-result paper of
479 this scope, with the specific main-body or audit-artifact evidence that addresses each. It is intended
480 as a reading guide rather than a substitute for the body.

481 **O1: “The covariance budget is too narrow; a richer posterior would pass.”** Section 4 reports
482 a rank-16/32/64/128 low-rank Hessian-plus-diagonal Laplace sweep, five exact block-diagonal and
483 joint-group full-covariance rows covering 22,064 to 270,896 trainable weights, and a streamed exact
484 joint-group probe over the full 270,896-weight vector. The rank-128 row is the only configuration
485 that crosses any gate axis (mask-overlap 0.816 with logit/activation CKA 0.932/0.910), and even
486 there layer-sparsity KS and basin-count axes still fail. The exact full-network full-covariance literal
487 baseline is bounded by the 553.1 GiB / 1106.3 GiB feasibility audit, not silently dropped.

488 **O2: “The result is a BatchNorm running-statistics artifact.”** The full CIFAR sweep over
489 freeze, recalibrate, and dense_buffers posterior BN policies, for both SGLD and cyclical
490 SGLD, has been executed and integrated. The eleven planned rerun entries are all observed (one
491 validation-selected, one locked final-test, six BN policy ablations, three saved-artifact reruns). The
492 single configuration that collapses—vanilla SGLD with frozen BN—is reported as an invalid sam-
493 pler setting rather than evidence for the posterior-mode account, and the remaining five of six BN
494 rows preserve the negative direction.

495 **O3: “Channel alignment is too weak; the posterior matches IMP after a permutation.”**
496 Activation-channel and weight-correlation alignment, record-level post-hoc matching, a block-
497 coordinate global-channel audit on the saved full-data artifact, and an exact stage-1 enumeration
498 over a small ResNet subgraph are all reported in the main body. The block-coordinate global-
499 channel audit leaves the global posterior/ticket Hamming near 0.21 (well below the 0.7 threshold),
500 and exhaustive full-data graph isomorphism is bounded as infeasible rather than silently skipped.

501 **O4: “This is only CIFAR-10 ResNet-20; the result will flip under another architecture or
502 dataset.”** The TinyCNN architecture-general cell in Section 4 runs the same matched five-seed
503 direct probe on a non-residual three-convolution CNN with no global average pooling, and repro-
504 duces the seed-level pattern: IMP beats the dense chain start 5/5, posterior samples stay below both
505 5/5, and the posterior-to-chain-start mask Hamming is 0.061 (95% CI [0.060, 0.062]). The dataset
506 axis is now confirmed under the validation-selected CIFAR-100 locked test (Section 4); SVHN and
507 larger natural-image benchmarks remain scoped limitations in Section 7.

508 **O5: “A negative result alone does not justify publication.”** The contribution is intentionally
509 three-part: (i) the controlled negative empirical result under a *pre-specified* support-equivalence
510 gate, with locked-test pre-registered predictions all confirmed (Section 3); (ii) the positive trajectory
511 and pruning-process account with Proposition 1 and the IMP-only residual mechanism (Section 5);
512 and (iii) a reusable posterior-vs-IMP audit framework (operational gate, fourteen-approximation
513 harness, pre-registered locked-test protocol, and audit suite) that any future posterior family can be
514 tested against.

515 B Representative Movement Diagnostics

516 C Generated Evidence Tables

Sampler	LR/Scale/Step	Dense	IMP	Post.	Chain	Post-chain	Rewind
SGLD	10^{-6}	0.8845	0.8993	0.1425	0.1441	0.7362	0.1784
SGLD	3×10^{-6}	0.8845	0.8993	0.1381	0.1441	0.5928	0.1784
cSGLD	10^{-6}	0.8852	0.8960	0.1422	0.1454	0.7046	0.1789
DiagLap	10^{-3}	0.8849	0.8980	0.1447	0.1469	0.8826	0.1787
DiagLap	10^{-2}	0.8849	0.8980	0.1278	0.1469	0.5961	0.1787
KFACLap	10^{-2}	0.8860	0.8956	0.1303	0.1456	0.4859	0.1775
LowRank32Lap	10^{-2}	0.8853	0.8964	0.1358	0.1457	0.7402	0.1795
LowRank64Lap	10^{-2}	0.8874	0.8954	0.1339	0.1433	0.7397	0.1766
LowRank128Lap	10^{-2}	0.8867	0.8967	0.1351	0.1453	0.7358	0.1780
SWAG20	64	0.8859	0.8973	0.1453	0.1455	0.9086	0.1782
RandSubHMC	3×10^{-3}	0.8866	0.8965	0.1440	0.1440	0.9766	0.1779
TrajSubHMC	10^{-3}	0.8849	0.8961	0.2290	0.2292	0.9915	0.1792
Hess16SubHMC	3×10^{-4}	0.8873	0.8974	0.1468	0.1468	0.9994	0.1799
Hess32SubHMC	3×10^{-4}	0.8868	0.8973	0.1461	0.1461	0.9993	0.1788

Table 2: Representative five-seed CIFAR-10 ResNet-20 epoch-1 rewind movement diagnostics at r5 p0.30. Posterior support either remains near the chain-start magnitude support or moves away from it, but posterior-to-IMP overlap (Post.) does not exceed the matching chain-start control. Traj-SubHMC is higher because its chain-start control is the matched dense trajectory magnitude support; the full generated movement summary is in Appendix Table 3.

Sampler	Scale	Post-chain	Post.-Chain	Rewind-Post.	Sample Acc.
DiagLap	10^{-3}	0.8826	-0.0021	0.0340	0.8799
KFACLap	10^{-3}	0.8016	-0.0016	0.0334	0.8839
KFACLap	10^{-2}	0.4859	-0.0154	0.0473	0.8695
LowRank32Lap	10^{-3}	0.9299	-0.0007	0.0345	0.8843
LowRank32Lap	10^{-2}	0.7402	-0.0098	0.0436	0.8789
LowRankLap	10^{-3}	0.9295	-0.0008	0.0330	0.8848
LowRankLap	10^{-2}	0.7359	-0.0105	0.0427	0.8784
SGHMC	10^{-7}	0.6796	-0.0037	0.0358	0.8752
SGLD	10^{-6}	0.7362	-0.0017	0.0359	0.8753
SWAG20	1.6×10^1	0.9528	-7.0e-05	0.0328	0.8636
SWAG20	6.4×10^1	0.9086	-0.0002	0.0329	0.8041
cSGLD	10^{-6}	0.7046	-0.0033	0.0368	0.8782

Table 3: Paired five-seed CIFAR movement statistics. Post.-Chain is posterior-to-IMP minus chain-start-to-IMP; Rewind-Post. is rewind-magnitude-to-IMP minus posterior-to-IMP.

Sampler	Scale	Head post-chain	Head post.-chain	Head rewind-post.	Sample Acc.
HeadLap	10^{-3}	0.7773	-0.0085	0.0209	0.8856
HeadLap	10^{-2}	0.6912	-0.0284	0.0407	0.8855
HeadLap	10^0	0.6769	-0.0352	0.0475	0.8835

Table 4: Exact full-covariance Laplace probe for the final CIFAR classification head only. Head post.-chain is head posterior-to-IMP minus head chain-start-to-IMP; head rewind-post. is head rewind-magnitude-to-IMP minus head posterior-to-IMP.

Sampler	Block	Params	B post-chain	B post.-chain	B rewind-post.	G post.-chain	Acc.
BlockDiagLap	11-block diag	22064	0.5158	-0.0114	0.0667	0.0036	0.8810
BlockDiagLap	16-block diag	68144	0.6146	-0.0050	0.0520	0.0010	0.8802
BlockLap	layer1.0.conv1	2304	0.2236	-0.0075	0.0464	0.0013	0.8961
BlockLap	layer3.0.shortcut.0	2048	0.3626	-0.0010	0.0648	0.0008	0.8905
JointDiagLap	joint-group diag 20k	86576	0.7009	-0.0023	0.0465	0.0006	0.8828
JointDiagLap	joint-group diag 10k	68144	0.5804	-0.0050	0.0523	0.0015	0.8811
JointDiagLap	joint-group diag 40k	270896	0.7389	-0.0019	0.0362	-0.0019	0.8824
JointLap	joint group	5424	0.5088	-0.0206	0.0343	0.0010	0.8922

Table 5: Full-covariance block Laplace probes for selected CIFAR ResNet-20 weight tensors, tensor groups, or a tensor-block diagonal subset with the rest of the network frozen. Block post.-chain is block posterior-to-IMP minus block chain-start-to-IMP; global post.-chain reports the induced global support change.

Sampler	Step	PostCh.	Post.-Chain	Rwd.-Post.	Acc.	Accept	ParamDist.
Hess16SubHMC	3×10^{-4}	0.9994	-1.6e-05	0.0331	0.8881	0.8833	0.0095
Hess32SubHMC	3×10^{-4}	0.9993	3.6e-05	0.0326	0.8872	0.9500	0.0104
HessSubHMC	3×10^{-4}	0.9999	2.3e-06	0.0339	0.8865	0.8600	0.0022
RandSubHMC	3×10^{-3}	0.9766	-2.2e-05	0.0339	0.8863	0.7400	0.3672
TrajSubHMC	3×10^{-4}	0.9950	-0.0001	-0.0499	0.8849	0.8000	0.0803
TrajSubHMC	10^{-3}	0.9915	-0.0002	-0.0498	0.8847	0.6900	0.1589

Table 6: Full-network low-dimensional subspace HMC probes around the dense CIFAR checkpoint. RandSubHMC and TrajSubHMC use random and dense-trajectory subspaces; Hessian rows use randomized top-Hessian subspaces.

Row	Post.	Chain	Delta	Wins	Post-chain	W-dist.
JointLap	0.3294	0.3501	-0.0206	0.0000	0.5088	0.0206
DiagLap	0.1447	0.1469	-0.0021	0.0000	0.8826	0.0021
HeadLap	0.6983	0.7068	-0.0085	0.4000	0.7773	0.0125
Hess16HMC	0.1468	0.1468	-1.6e-05	0.4000	0.9994	3.8e-05
Hess32HMC	0.1461	0.1461	3.6e-05	0.8000	0.9993	3.7e-05
KFACLap	0.1456	0.1456	-2.2e-05	0.6000	0.9334	0.0002
LowRank32Lap	0.1358	0.1457	-0.0098	0.0000	0.7402	0.0098
LowRankLap	0.1351	0.1456	-0.0105	0.0000	0.7359	0.0105
RandHMC	0.1440	0.1440	-2.2e-05	0.2000	0.9766	6.8e-05
SGHMC move	0.1419	0.1457	-0.0037	0.0000	0.6796	0.0037
SGLD move	0.1425	0.1441	-0.0017	0.0000	0.7362	0.0017
SGLD-3	0.1368	0.1368	-4.9e-06	0.4000	0.9969	3.4e-05
SWAG	0.1361	0.1361	-5.4e-05	0.6000	0.9265	0.0003
SWAG20	0.1454	0.1455	-7.0e-05	0.2000	0.9528	0.0001
TrajHMC	0.2290	0.2292	-0.0002	0.4000	0.9915	0.0002
cSGLD move	0.1422	0.1454	-0.0033	0.0000	0.7046	0.0033
Fashion	0.2114	0.2122	-0.0008	0.0000	0.9212	0.0008
MNIST	0.2373	0.2385	-0.0011	0.0000	0.9298	0.0011

Table 7: Mode/ticket distribution-equivalence audit for selected posterior families. Delta is posterior-to-IMP minus matched chain-start-to-IMP. Wins is the fraction of seed-level paired posterior-chain deltas above zero; W-dist. is the two-sample Wasserstein distance over grouped seed means. A positive posterior mode account would require positive deltas rather than practical ties to chain-start magnitude supports.

Setting	Comparison	Modes	H_n	n_L	n_T	KS p	Ham.	Logit	Act.	H-cost	A-cost	Pass
Digits MLP	posterior samples vs tickets	1	0.0000	50	5	0.0788	0.6314	0.9819	n/a	0.0181	n/a	2/4
Digits MLP	posterior modes vs tickets	1	0.0000	1	5	0.9216	n/a	0.9821	n/a	0.0179	n/a	3/4
CIFAR subset ResNet	posterior samples vs tickets	1	0.0000	25	5	1.1e-05	0.6000	0.8878	n/a	0.1122	n/a	2/4
CIFAR subset ResNet	posterior modes vs tickets	1	0.0000	1	5	0.1388	n/a	0.9077	n/a	0.0923	n/a	3/4
CIFAR subset act. pilot	posterior samples vs tickets	3	0.4057	15	3	0.0330	0.6571	0.9022	0.8726	0.0978	0.1274	4/6
CIFAR subset act. pilot	posterior modes vs tickets	3	0.4057	3	3	0.2264	1.0000	0.9014	0.8704	0.0986	0.1296	6/6
CIFAR full ResNet	posterior samples vs tickets	1	0.0000	50	5	5.3e-09	0.0033	0.9369	0.9172	0.0631	0.0828	4/6
CIFAR full ResNet	posterior modes vs tickets	1	0.0000	1	5	0.0329	n/a	0.9383	0.9179	0.0617	0.0821	4/6
CIFAR full aligned	aligned posterior samples vs tickets	1	0.0000	50	5	2.3e-09	0.0000	0.9373	0.9168	0.0627	0.0832	4/6
CIFAR full aligned	aligned posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9363	0.9160	0.0637	0.0840	4/6
CIFAR full weight-aligned	weight-aligned chain start magnitude vs tickets	1	0.0000	5	5	4.8e-05	0.0000	0.9374	0.9195	0.0626	0.0805	4/6
CIFAR full weight-aligned	weight-aligned posterior samples vs tickets	1	0.0000	50	5	1.2e-08	0.1290	0.9336	0.9131	0.0664	0.0869	4/6
CIFAR full weight-aligned	weight-aligned posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9327	0.9132	0.0673	0.0868	4/6
CIFAR full cSGLD multi-chain	chain start magnitude vs tickets	1	0.0000	15	5	5.2e-07	0.0857	0.9363	0.9196	0.0637	0.0804	4/6
CIFAR full cSGLD multi-chain	posterior samples vs tickets	1	0.0000	75	5	3.3e-08	0.2461	0.9327	0.9144	0.0673	0.0856	4/6
CIFAR full cSGLD multi-chain	posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9288	0.9115	0.0712	0.0885	4/6
CIFAR full cSGLD independent	chain start magnitude vs tickets	1	0.0000	15	5	7.6e-08	0.0000	0.9306	0.9105	0.0694	0.0895	4/6
CIFAR full cSGLD independent	posterior samples vs tickets	1	0.0000	75	5	9.3e-10	0.0000	0.9269	0.9051	0.0731	0.0949	4/6
CIFAR full cSGLD independent	posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9242	0.9054	0.0758	0.0946	4/6
CIFAR full LowRank128Lap	chain start magnitude vs tickets	1	0.0000	5	5	9.0e-05	0.0000	0.9385	0.9218	0.0615	0.0782	4/6
CIFAR full LowRank128Lap	posterior samples vs tickets	1	0.0000	50	5	2.0e-06	0.8163	0.9319	0.9096	0.0681	0.0904	5/6
CIFAR full LowRank128Lap	posterior modes vs tickets	1	0.0000	1	5	0.1388	n/a	0.9289	0.9023	0.0711	0.0977	5/6
CIFAR full JointDiagLap270k	chain start magnitude vs tickets	1	0.0000	5	5	2.5e-05	0.0000	0.9389	0.9221	0.0611	0.0779	4/6
CIFAR full JointDiagLap270k	posterior samples vs tickets	1	0.0000	25	5	1.1e-08	0.0000	0.9373	0.9199	0.0627	0.0801	4/6
CIFAR full JointDiagLap270k	posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9379	0.9225	0.0621	0.0775	4/6

Table 8: Direct proposal-level mode/ticket distribution probe on digits MLP plus CIFAR-10 ResNet-20 subset and full-data settings. Posterior sample masks and mean-shift posterior-mode representations are compared with IMP tickets using the proposal’s layer-sparsity KS, mask-Hamming overlap, logit-space CKA, final-hidden-activation CKA, and Hungarian-cost thresholds; Modes and H_n report mean-shift cluster count and normalized raw-parameter basin entropy.

Source	Acc.	NLL	ECE	MSP AUROC	Ent. AUROC	MSP FPR95
Dense	0.8866	0.3536	0.0353	0.8230	0.8326	0.7636
IMP	0.8953	0.3387	0.0393	0.8306	0.8392	0.7508
SWAG ens.	0.8688	0.4018	0.0285	0.8050	0.8171	0.7715
Learned random	0.8449	0.4632	0.0270	0.7897	0.8024	0.7870
Gem-Miner	0.8418	0.4694	0.0283	0.7853	0.7985	0.7969
Var. prune	0.8301	0.5008	0.0255	0.7754	0.7897	0.8101

Table 9: Five-seed CIFAR-10 calibration and CIFAR-100 OOD diagnostics. SWAG ens. averages ten SWAG predictive samples per seed. Learned random, Gem-Miner, and Var. prune rows are trained from initialization at the IMP sparsity. MSP and entropy AUROC use maximum softmax probability and negative predictive entropy as in-distribution scores; lower FPR95 is better.

Source	Acc.	Support-IMP	Support-Dense	Support-Rewind
Epoch 1	0.4017	0.1782	0.2231	1.0000
Epoch 5	0.6881	0.2197	0.3582	0.3724
Epoch 10	0.7630	0.2342	0.5350	0.2746
Epoch 20	0.8492	0.2328	0.8433	0.2273
Epoch 30	0.8841	0.2312	1.0000	0.2231
max abs	—	0.2342	0.5616	0.3517
mean abs	—	0.2388	0.5314	0.4214
rms abs	—	0.2400	0.5713	0.4046

Table 10: Five-seed CIFAR dense-trajectory controls. The epoch-1 state is the IMP rewind point. Checkpoint rows use magnitude supports at that epoch; aggregate rows use trajectory score masks across checkpoints.

Source	Kind	Accuracy	Acc.-IMP	Support-IMP
IMP	IMP	0.8983	0.0006	1.0000
random 0	random	0.8422	-0.0555	0.0922
Epoch 1	checkpoint	0.8540	-0.0437	0.1785
Epoch 10	checkpoint	0.8730	-0.0247	0.2374
Epoch 30	checkpoint	0.8826	-0.0151	0.2356
rms abs	aggregate	0.8739	-0.0238	0.2429
mean abs	aggregate	0.8743	-0.0234	0.2421
path length	aggregate	0.8549	-0.0428	0.1772
rewind rms movement	aggregate	0.8572	-0.0405	0.1639
Gem-Miner	Gem-Miner	0.8471	-0.0499	0.0917
Var. prune	Var. prune	0.8306	-0.0669	0.0907
Hard concrete	Hard concrete	0.2766	-0.6204	0.0922

Table 11: Five-seed CIFAR trajectory-mask retraining controls. All masks are trained for 30 epochs from the same epoch-1 rewind state at the IMP sparsity. Acc.-IMP is candidate accuracy minus the final IMP accuracy from the corresponding seed.

Source	Kind	Acc.	ECE	Brier	Acc.-IMP	Supp.-IMP	Supp.-Dense
IMP	IMP	0.9711	0.0211	0.0410	0.0000	1.0000	0.7604
random 0	random	0.9489	0.0398	0.0821	-0.0222	0.2151	0.2139
Gem-Miner	Gem-Miner	0.9550	0.0392	0.0753	-0.0161	0.2081	0.2082
Var. prune	Var. prune	0.9633	0.0250	0.0532	-0.0078	0.3822	0.4010

Table 12: Five-seed digits MLP variational pruning sanity check. Var. prune optimizes Bernoulli mask probabilities on frozen initialization weights with Concrete samples, KL, sparsity, and entropy penalties before retraining the selected hard mask.

Base	Variant	Alpha	Accuracy	Acc.-IMP
Epoch 30	top IMP	0	0.8797	-0.0175
Epoch 30	top IMP	0.5	0.8882	-0.0090
Epoch 30	random non-IMP	0.5	0.8780	-0.0192
Epoch 30	top IMP	1	0.8966	-0.0006
rms abs	top IMP	0	0.8733	-0.0238
rms abs	top IMP	0.5	0.8851	-0.0121
rms abs	random non-IMP	0.5	0.8705	-0.0266
Epoch 10	top IMP	0	0.8712	-0.0259
Epoch 10	top IMP	0.5	0.8855	-0.0116
Epoch 10	random non-IMP	0.5	0.8704	-0.0268

Table 13: Five-seed CIFAR trajectory residual controls. Starting from a base trajectory mask, an IMP-residual row swaps the indicated fraction of base-only support for IMP-only support; a random-residual row swaps the same number of positions for non-IMP random support.

Base	Base	Low rm	Rand. rm	High rm	Rand. non
Epoch 30	0.8792	0.8881	0.8883	0.8906	0.8779
rms abs	0.8738	0.8874	0.8896	0.8914	0.8709
Epoch 10	0.8687	0.8862	0.8922	0.8920	0.8701

Table 14: Five-seed CIFAR residual removal-order controls. Low/Rand./High rm add the same top IMP-only residual support and vary which base-only weights are removed by low, random, or high base score. Rand. non removes low base-score weights but adds same-size non-IMP random support.

Base	Support-IMP	IMP-only(k)	Prune rnd.	Pred. AUC	Pred. recall
Epoch 30	0.2323	28.4k	2.9003	0.6197	0.2175
rms abs	0.2412	27.8k	2.9745	0.6165	0.2087
Epoch 10	0.2359	28.1k	2.9669	0.6206	0.2206

Table 15: Five-seed CIFAR residual anatomy. IMP-only(k) is the number of IMP-kept weights missing from the base support. Prune rnd. is the mean IMP pruning round for base-only weights. The predictor is a held-out logistic model over dense-trajectory rank features plus stage indicators, trained to recover IMP-only residual support among non-base weights.

Base	Base	Oracle	Pred.	Random	Pred. prec.	Rand. prec.
Epoch 30	0.8808	0.8892	0.8793	0.8805	0.1866	0.1253
rms abs	0.8740	0.8866	0.8744	0.8744	0.1834	0.1233
Epoch 10	0.8712	0.8890	0.8728	0.8724	0.1843	0.1239

Table 16: Five-seed CIFAR functional residual-predictor masks. For each base, the held-out predictor swaps half of base-only support for non-base weights with highest predicted IMP-only probability. Oracle uses true IMP-only residual support, and random uses same-size held-out non-IMP-random support. Prec. is the IMP-only precision of added weights.

Base	Base	Oracle	Cross-seed	Random	Cross prec.	Rand. prec.
Epoch 30	0.8815	0.8905	0.8776	0.8781	0.2413	0.1260
rms abs	0.8733	0.8878	0.8731	0.8725	0.2238	0.1246
Epoch 10	0.8730	0.8890	0.8745	0.8726	0.2388	0.1264

Table 17: Five-seed CIFAR cross-seed residual-transfer masks. For each held-out target seed, the predictor is trained on non-base candidate weights from the other four seeds, then swaps half of target base-only support for target non-base weights with highest predicted IMP-only probability. Oracle and random use the target seed only.

Base	Base	Oracle	Src vote	Aln src	Vote rand.	Aln rand.	Tgt rand.	Src prec.
Epoch 30	0.8813	0.8886	0.8769	0.8797	0.8774	0.8809	0.8779	0.1513
rms abs	0.8737	0.8872	0.8739	0.8727	0.8735	0.8743	0.8728	0.1451
Epoch 10	0.8720	0.8890	0.8725	0.8710	0.8732	0.8740	0.8719	0.1491

Table 18: Five-seed CIFAR direct cross-seed residual-support transfer. Source-vote variants add target non-base coordinates selected by votes from the other seeds’ oracle residual supports. Aln src and Aln rand. first map source channels to target channels by activation-correlation Hungarian matching. Tgt rand. is a target-seed random non-base residual control. Src prec. is target IMP-only precision among source-vote additions.

Base	T base	T oracle	T rand.	M base	M oracle	M rand.
Epoch 30	0.8827	0.8893	0.8791	0.8641	0.8926	0.8649
rms abs	0.8743	0.8892	0.8745	0.8605	0.8910	0.8628
Epoch 10	0.8744	0.8892	0.8712	0.8607	0.8942	0.8636

Table 19: Five-seed CIFAR residual base-compatibility controls. T rows use the actual trajectory base. M rows replace that base with a per-parameter random base preserving the same IMP/non-IMP counts and therefore the same base-to-IMP overlap before adding oracle or random residual support.

Base	M base	Top	Dense	RMS	RMS-d	Std	Rand	Dense ovlp	RMS ovlp	Rand ovlp
Epoch 30	0.8606	0.8915	0.8821	0.8852	0.8770	0.8710	0.8783	0.5573	0.5560	0.5006
rms abs	0.8631	0.8928	0.8812	0.8834	0.8745	0.8717	0.8795	0.5528	0.5510	0.5011
Epoch 10	0.8609	0.8911	0.8827	0.8829	0.8764	0.8710	0.8791	0.5565	0.5556	0.5002

Table 20: Five-seed CIFAR residual posterior-decomposition controls on IMP-overlap-matched random bases. Top adds the target oracle IMP-only residual. Dense, RMS, RMS-d, Std, and Rand add the same number of target IMP-only residual weights ranked by dense final magnitude, diagonal-Laplace posterior RMS, posterior RMS minus dense magnitude, posterior standard deviation, or uniform sampling. Ovlp columns report the fraction of additions that also belong to the top oracle additions.

Base	Base	Oracle	Rand.	IMP	Global non	Param+score non
Epoch 30	0.8803	0.8872		0.8818	0.8767	0.8758
rms abs	0.8708	0.8858		0.8794	0.8701	0.8716
Epoch 10	0.8678	0.8854		0.8764	0.8695	0.8685

Table 21: Five-seed CIFAR residual stratified controls. All non-base variants remove the same low-base-score weights as the oracle. Rand. IMP adds random IMP-only residual weights. Global non adds random non-IMP non-base weights. Param+score non matches the oracle added weights by parameter tensor and within-parameter score decile, but excludes IMP-only weights.

Base	Base	Oracle	R1 acc.	R3 acc.	R5 acc.	R1 prec.	R3 prec.
Epoch 30	0.8793	0.8898	0.8792	0.8829	0.8867	0.4431	0.7523
rms abs	0.8736	0.8881	0.8771	0.8832	0.8857	0.4308	0.7459
Epoch 10	0.8724	0.8884	0.8772	0.8826	0.8861	0.4395	0.7536

Table 22: Five-seed CIFAR residual IMP-process probe. R1/R3/R5 rows add half residual support from IMP round-survivor candidates ranked by the corresponding round-trained weights. Prec. is the fraction of added weights that survive in the final IMP mask; R5 precision is one by construction.

Base	Rnd	Top acc.	Rand. acc.	Low acc.	Top final	Rand. final	Low final	Top oracle	Rand. oracle	Low oracle
Epoch 30	1	0.8806	0.8774	0.8776	0.4373	0.1913	0.0811	0.2706	0.0960	0.0340
Epoch 30	3	0.8839	0.8813	0.8770	0.7530	0.4357	0.1881	0.4777	0.2176	0.0652
Epoch 30	5	0.8887	0.8847	0.8815	1.0000	1.0000	1.0000	0.6818	0.4998	0.3182
rms abs	1	0.8788	0.8733	0.8738	0.4268	0.1871	0.0808	0.2617	0.0924	0.0343
rms abs	3	0.8830	0.8759	0.8726	0.7496	0.4294	0.1885	0.4734	0.2152	0.0657
rms abs	5	0.8857	0.8803	0.8759	1.0000	1.0000	1.0000	0.6803	0.4987	0.3197
Epoch 10	1	0.8795	0.8724	0.8699	0.4358	0.1880	0.0818	0.2698	0.0933	0.0345
Epoch 10	3	0.8822	0.8746	0.8727	0.7555	0.4347	0.1898	0.4795	0.2178	0.0657
Epoch 10	5	0.8855	0.8805	0.8755	1.0000	1.0000	1.0000	0.6821	0.4992	0.3179

Table 23: Five-seed CIFAR residual IMP-process ranking controls. Top uses the highest round-trained scores among round-survivor candidates; Rand. samples the same round-survivor candidate set; Low uses the lowest round-trained scores. Final is the fraction of added weights that survive in the final IMP mask; oracle is overlap with the final oracle added subset.

Base	Rnd	Score acc.	Match acc.	Δ	Wins	Oracle
Epoch 10	1	0.8823	0.8799	0.0024	5/5	0.5568
Epoch 10	3	0.8825	0.8789	0.0036	4/5	0.5976
Epoch 10	5	0.8851	0.8820	0.0030	5/5	0.6792
Epoch 30	1	0.8841	0.8843	-0.0002	3/5	0.5569
Epoch 30	3	0.8842	0.8831	0.0011	4/5	0.5971
Epoch 30	5	0.8860	0.8845	0.0014	3/5	0.6788
rms abs	1	0.8824	0.8787	0.0038	4/5	0.5538
rms abs	3	0.8821	0.8808	0.0013	3/5	0.5943
rms abs	5	0.8854	0.8836	0.0018	4/5	0.6773

Table 24: Five-seed CIFAR residual IMP-process oracle-overlap-matched control. Score acc. uses the round-trained score ordering inside the final-IMP residual candidate set. Match acc. samples random final-IMP residual additions with the same final-oracle overlap count. Δ is score minus matched random accuracy; wins are positive paired seed deltas.

Base	Rnd	Round acc.	Dense acc.	Base acc.	Δ_D	WinD	Δ_B	WinB	Oracle
Epoch 10	1	0.8833	0.8827	0.8824	0.0005	3/5	0.0008	4/5	0.5559
Epoch 10	3	0.8838	0.8827	0.8824	0.0011	4/5	0.0014	4/5	0.5984
Epoch 10	5	0.8872	0.8827	0.8824	0.0045	5/5	0.0048	5/5	0.6780
Epoch 30	1	0.8861	0.8831	0.8831	0.0030	5/5	0.0030	5/5	0.5574
Epoch 30	3	0.8874	0.8831	0.8831	0.0042	4/5	0.0042	4/5	0.5992
Epoch 30	5	0.8873	0.8831	0.8831	0.0041	5/5	0.0041	5/5	0.6779
rms abs	1	0.8820	0.8823	0.8820	-0.0004	2/5	-4.0e-05	2/5	0.5536
rms abs	3	0.8844	0.8823	0.8820	0.0021	4/5	0.0024	5/5	0.5950
rms abs	5	0.8870	0.8823	0.8820	0.0046	5/5	0.0050	5/5	0.6760

Table 25: Five-seed CIFAR residual IMP-process score-source controls. Candidate set is final-IMP residual for all variants; dense/base controls rank that candidate set by dense-final or base-source magnitudes. Δ_D/Δ_B are round-score minus dense/base score accuracy; wins are positive paired seed deltas.

Base	Rnd	Round acc.	Excl. acc.	Δ	Wins	Round oracle	Excl. oracle
Epoch 10	1	0.8844	0.8770	0.0074	5/5	0.5602	0.4398
Epoch 10	3	0.8858	0.8782	0.0076	5/5	0.5995	0.4005
Epoch 10	5	0.8872	0.8779	0.0093	5/5	0.6762	0.3238
Epoch 30	1	0.8848	0.8830	0.0018	4/5	0.5608	0.4392
Epoch 30	3	0.8867	0.8834	0.0033	5/5	0.5983	0.4017
Epoch 30	5	0.8867	0.8824	0.0043	5/5	0.6755	0.3245
rms abs	1	0.8846	0.8779	0.0067	5/5	0.5581	0.4419
rms abs	3	0.8858	0.8777	0.0081	5/5	0.5959	0.4041
rms abs	5	0.8861	0.8793	0.0068	5/5	0.6733	0.3267

Table 26: Five-seed CIFAR residual IMP-process round-exclusion control. Round acc. uses the round-trained score ordering inside the final-IMP residual candidate set. Excl. acc. removes the round-selected additions from that candidate set, then chooses the best remaining final-IMP residual additions by final IMP magnitude under the same support budget. Δ is round minus excluded accuracy; wins are positive paired seed deltas.

Base	Rnd	Round acc.	Excl. acc.	Tensor excl.	Δ_E	WinE	Δ_T	WinT	Tensor oracle
rms abs	5	0.8855	0.8753	0.8764	0.0103	5/5	0.0091	5/5	0.4153

Table 27: Five-seed CIFAR residual IMP-process tensor-matched round-exclusion control at the RMS trajectory base and round 5. Tensor excl. removes the round-selected additions, then chooses replacement final-IMP residual additions matched to the removed coordinates by parameter tensor before applying the final IMP magnitude score. Δ_E is round minus the unrestricted excluded oracle; Δ_T is round minus the tensor-matched excluded oracle.

Base	Rnd	Round acc.	Tensor excl.	Tensor+score excl.	Δ_T	WinT	Δ_{TS}	WinTS
rms abs	5	0.8878	0.8784	0.8837	0.0095	5/5	0.0041	5/5

Table 28: Five-seed CIFAR residual IMP-process tensor+score-matched round-exclusion control at the RMS trajectory base and round 5. Tensor+score excl. removes the round-selected additions, then chooses replacement final-IMP residual additions matched to the removed coordinates by parameter tensor and within-tensor round-score decile before applying final IMP magnitude. Δ_T is round minus the tensor-matched excluded oracle; Δ_{TS} is round minus the tensor+score-matched excluded oracle.

Base	Rnd	Round acc.	Resid. acc.	Dense acc.	Base acc.	Δ_R	WinR	Oracle drop
rms abs	5	0.8852	0.8811	0.8814	0.8811	0.0041	5/5	0.1830

Table 29: Five-seed CIFAR residual IMP-process projection control at the RMS trajectory base and round 5. Resid. acc. uses a round-score ordering after linearly projecting out the base-source, dense-final, and final-IMP magnitude scores inside the final-IMP residual candidate pool. Δ_R is round minus residualized-score accuracy; Oracle drop is round oracle overlap minus residualized oracle overlap.

Base	Rnd	Round acc.	Post-resid. acc.	Dense acc.	Base acc.	Δ_P	WinP	Oracle drop
rms abs	5	0.8847	0.8825	0.8815	0.8797	0.0023	5/5	0.1923

Table 30: Five-seed CIFAR residual IMP-process posterior-projection control at the RMS trajectory base and round 5. Post-resid. acc. uses a round-score ordering after linearly projecting out the base-source, dense-final, final-IMP magnitude, and diagonal-Laplace posterior RMS/std/excess scores inside the final-IMP residual candidate pool. Δ_P is round minus posterior-residualized accuracy; Oracle drop is round oracle overlap minus posterior-residualized oracle overlap.

Base	Rnd	Round acc.	Subspace-resid. acc.	Dense acc.	Base acc.	Δ_S	WinS	Oracle drop
rms abs	5	0.8869	0.8821	0.8822	0.8822	0.0048	5/5	0.1890

Table 31: Five-seed CIFAR residual IMP-process learned-subspace control at the RMS trajectory base and round 5. Subspace-resid. acc. uses a round-score ordering after projecting out the top rank-8 PCA component scores learned from dense trajectory, final-IMP magnitude, and earlier IMP-round scores inside the final-IMP residual candidate pool. Δ_S is round minus learned-subspace-residualized accuracy; Oracle drop is round oracle overlap minus learned-subspace-residualized oracle overlap.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes].

Justification: The abstract and Introduction state the support-equivalence claim, the negative conclusion, and the exact dense full-network CIFAR covariance limitation.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes].

Justification: The Limitations and Next Experiments section lists the remaining posterior, permutation, learned-mask, process-causality, and reproducibility limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete proof?

Answer: [N/A].

Justification: The paper is empirical and does not claim new theoretical results.

4. Experimental result reproducibility

Question: Does the paper fully disclose the information needed to reproduce the main experimental results?

Answer: [Yes].

Justification: The main text and generated appendix summarize the settings, and the accompanying artifact package contains scripts, run summaries, environment locks, a Docker artifact verifier, and a release manifest.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to reproduce the main experimental results?

Answer: [Yes].

Justification: The anonymized source repository is public at <https://github.com/evaldataset/lottery-ticket-bayesian-modes-artifact>, and the public release page hosts the artifact archive with SHA256 sidecar at <https://github.com/evaldataset/lottery-ticket-bayesian-modes-artifact/releases>. The local receipt registry records the exact archive URL, source commit, and passing public CI run; external GPU-container validation remains pending and is tracked separately as a release-readiness limitation.

6. Experimental setting/details

Question: Does the paper specify all training and test details necessary to understand the results?

Answer: [Yes].

Justification: The Operational Test and Current Results sections describe datasets, model families, posterior approximations, seeds, controls, and comparison metrics; the appendix and artifact manifest give generated tables and commands.

7. Experiment statistical significance

Question: Does the paper report error bars or other suitable information about statistical significance?

Answer: [Yes].

Justification: Main results are five-seed summaries with confidence intervals, paired win counts, grouped comparisons, and generated statistical tables.

565 **8. Experiments compute resources**

566 Question: For each experiment, does the paper provide sufficient information on compute
567 resources needed to reproduce the experiments?

568 Answer: [Yes].

569 Justification: The artifact package includes CPU and CUDA environment locks,
570 container definitions, exact reproduction command lines, dense-covariance mem-
571 ory bounds, artifact storage estimates, and compute-resource accounting in
572 docs/compute_resource_accounting.md. Historical metrics do not include profiler-
573 grade wall-clock fields, and that limitation is documented there.

574 **9. Code of ethics**

575 Question: Does the research conform with the NeurIPS Code of Ethics?

576 Answer: [Yes].

577 Justification: The work uses standard public benchmark datasets and synthetic/fake-data
578 smoke tests, does not involve human subjects, and reports negative results and limitations
579 transparently.

580 **10. Broader impacts**

581 Question: Does the paper discuss both potential positive and negative societal impacts of
582 the work?

583 Answer: [N/A].

584 Justification: This is foundational empirical analysis of pruning and posterior-support ex-
585 planations, with no direct deployed system or new capability; the main impact is method-
586 ological transparency around negative results.

587 **11. Safeguards**

588 Question: Does the paper describe safeguards for responsible release of data or models that
589 have a high risk for misuse?

590 Answer: [N/A].

591 Justification: The work does not release high-risk models, scraped datasets, generative
592 systems, or systems intended for deployment.

593 **12. Licenses for existing assets**

594 Question: Are existing assets properly credited and are license and terms of use explicitly
595 mentioned and respected?

596 Answer: [Yes].

597 Justification: Dataset, NeurIPS-style, and software dependency sources and license/terms
598 notes are inventoried in docs/asset_license_inventory.md. Raw benchmark datasets
599 are not redistributed in the local release manifest.

600 **13. New assets**

601 Question: Are new assets introduced in the paper well documented and provided alongside
602 the assets?

603 Answer: [Yes].

604 Justification: Project-created source, generated tables/figures, run summaries, saved
605 mask artifacts, reproducibility docs, container definitions, and excluded local
606 caches are documented in docs/new_asset_inventory.md and hash-inventoried by
607 docs/public_release_manifest.md.

608 **14. Crowdsourcing and research with human subjects**

609 Question: For crowdsourcing experiments and research with human subjects, does the pa-
610 per include the required details?

611 Answer: [N/A].

612 Justification: The paper does not use crowdsourcing or human-subject experiments.

613 **15. Institutional review board approvals**

614 Question: Does the paper describe IRB approvals or equivalent review for research with
615 human subjects?

616 Answer: [N/A].

617 Justification: The paper does not involve human-subject research.

618 **16. Declaration of LLM usage**

619 Question: Does the paper describe LLM usage if it is an important, original, or non-
620 standard component of the core methods?

621 Answer: [N/A].

622 Justification: LLMs are not part of the scientific method, experiments, posterior approxi-
623 mations, pruning algorithms, or analysis pipeline studied in the paper.