
Winning Tickets Are Not Posterior Modes: Evidence from Posterior Support and Magnitude Controls

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The lottery ticket hypothesis suggests that dense neural networks contain sparse
2 subnetworks that can be trained in isolation when rewound to their original ini-
3 tialization. A tempting Bayesian interpretation is that these tickets correspond
4 to posterior modes, or to posterior-basin supports. We turn this interpretation
5 into a pre-specified, falsifiable test: posterior-induced masks must match IMP
6 tickets better than local magnitude, dense-trajectory, and pruning-process con-
7 trols, not merely better than uniform random masks. Across five-seed MNIST and
8 Fashion-MNIST sparsity sweeps and a CIFAR-10 ResNet-20 epoch-1 rewind set-
9 ting, twelve posterior approximations—SGLD, SGHMC, cyclical SGLD, SWAG,
10 the Laplace family up to joint-group full covariance, and subspace HMC—fail this
11 stronger criterion. Posterior supports often beat random masks, but they remain
12 tied to chain-start magnitude, dominated by dense magnitude, or closer to rewind
13 controls than to IMP. Support-equivalence is graded rather than binary: only a
14 rank-128 Hessian-plus-diagonal Laplace crosses the mask-overlap threshold, and
15 even it fails the layer-sparsity and basin-count gates, so a posterior account sur-
16 vives on one axis at one fidelity. We then give a positive account of what winning
17 tickets are: trajectory and pruning-process objects. Dense-trajectory magnitude
18 explains most of the fixed-mask training gap and a process-selected IMP-only
19 residual explains the rest, with an extensive battery of process controls failing to
20 substitute for it. This account is consistent with axial-subspace views of IMP and
21 inconsistent with a posterior-mode view. Exact dense full-network full-covariance
22 CIFAR posterior baselines remain an open limitation.

1 Introduction

24 The lottery ticket hypothesis (LTH) states that a dense randomly initialized network contains sparse
25 subnetworks that can be trained to match the dense network when reset to their original initialization
26 [4]. This result raises a basic question about the object discovered by iterative magnitude pruning
27 (IMP): is a winning ticket a combinatorial artifact of the optimization trajectory, or does it reflect a
28 deeper probabilistic structure of the posterior over weights?

29 A Bayesian reading is attractive. SGD and its noisy variants have long been connected to approx-
30 imate Bayesian inference [16, 11]. Bayesian deep learning further emphasizes that useful neural
31 network solutions are distributed across broad posterior regions rather than isolated points [17, 7].
32 If winning tickets correspond to posterior modes or posterior-basin supports, LTH would be evi-
33 dence for a selection view of learning: pruning would reveal posterior structure already present in
34 the trained model.

35 This paper tests that interpretation directly. We ask:

36 Do posterior samples or posterior basins induce sparse supports that explain IMP
37 winning-ticket masks beyond dense magnitude, initialization magnitude, rewind
38 magnitude, chain-start magnitude, and pruning-process controls?

39 The paper makes three first-class contributions. **(i) A controlled negative empirical result.** We state
40 the posterior-mode explanation as a *pre-specified*, falsifiable support-equivalence claim—posterior-
41 induced supports must beat local magnitude controls, not just random masks, and mode/ticket distri-
42 butions must match in mask and function space—with the operational gates fixed before the poste-
43 rior runs. The gates are then evaluated across repeated MNIST/Fashion-MNIST sweeps and twelve
44 CIFAR-10 ResNet-20 posterior approximations (stochastic samplers, Laplace variants, SWAG, sub-
45 space HMC, and direct mode/ticket probes), with a locked final-test confirmation that matches three
46 pre-registered numerical predictions. Support-equivalence is *graded* rather than binary: a single
47 rank-128 Hessian-plus-diagonal Laplace passes the mask-overlap axis while still failing the layer-
48 sparsity and basin-count axes. **(ii) A positive trajectory and pruning-process account.** IMP tickets
49 are explained by dense-trajectory magnitude plus a process-selected IMP-only residual rather than
50 posterior-mode supports. Proposition 1 formalizes why high-SNR posterior TopK masks must track
51 the chain-start magnitude support. **(iii) A reusable posterior-vs-IMP audit framework.** The op-
52 erational support-equivalence gate, the twelve-approximation harness, the pre-registered locked-test
53 protocol, and the surrounding audit suite are released as a reusable testbed so other posterior families
54 can be tested against IMP tickets under the same operational contract.

55 The conclusion is negative but scoped. Posterior masks beat random masks, but they track the mag-
56 nitude support of the dense checkpoint or the chain start. When chains move, they move away
57 from the chain start without moving toward the IMP ticket. When function-space CKA is high,
58 mask-distribution and basin-count criteria still fail. When process controls are added, the strongest
59 positive signal is the IMP-selected residual, not a posterior-mode support. The cost is quantitative:
60 treating IMP tickets as posterior modes underperforms IMP by 2.3 percentage points on the locked
61 full-data CIFAR-10 ResNet-20 row (posterior 0.872 versus IMP 0.895) and by 5.3 percentage points
62 on the non-residual TinyCNN at matched sparsity (posterior 0.738 versus IMP 0.791), a direct ac-
63 curacy loss for posterior-based pruning pipelines.

64 2 Related Work

65 **Lottery tickets and pruning.** LTH was introduced by Frankle and Carbin [4]. Follow-up work
66 connected LTH to rewinding and linear mode connectivity [5]. Random weighted networks with
67 learned masks further emphasized the selection view of learning [13]. Single-shot pruning methods
68 such as SNIP [9] and SynFlow [15] provide important non-Bayesian controls.

69 **Mode connectivity and subspaces.** Low-loss connectivity between neural network solutions chal-
70 lenges the view that trained models are isolated modes [6, 2]. For lottery tickets specifically, Paul
71 et al. [12] argue that IMP masks encode axial subspaces that intersect matching linearly connected
72 regions, which is a direct non-Bayesian competitor to a posterior-mode explanation. Permutation
73 symmetries and re-basin methods further complicate parameter-space mode counting [3, 1]. Our ex-
74 periments therefore include function-space clustering and linear barrier probes in addition to mask
75 overlap. A central aim of this paper is to adjudicate between these two competitors directly: the
76 posterior-mode reading and the axial-subspace/trajectory reading make opposite predictions about
77 whether posterior-induced supports beat trajectory and pruning-process controls, and the operational
78 gates below are designed to separate them.

79 **Bayesian neural networks.** SGLD provides a simple route from SGD to approximate posterior
80 sampling [16]. SGD itself can be interpreted as approximate Bayesian inference in a stationary
81 regime [11]. Practical posterior approximations such as SWAG [10] and exact small-model sampling
82 studies [7] show that posterior geometry is structured and nontrivial. PAC-Bayesian and Bayesian
83 LTH analyses [14, 8] motivate the need for controls: Bayesian language alone does not imply that
84 IMP masks are posterior-mode supports.

85 Our contribution is therefore not a new posterior approximation or a claim that Bayesian inference
86 is uninformative for pruning. It is a controlled support-equivalence test: if posterior modes explain

winning tickets, their sparse supports should match IMP tickets after conditioning on dense magnitude, rewind magnitude, trajectory magnitude, and pruning-process alternatives.

3 Operational Test

Let m_{IMP} denote the mask found by IMP at a fixed sparsity. For a posterior sample $\theta^{(s)}$, the simplest posterior-induced mask is

$$m^{(s)} = \text{TopK}(|\theta^{(s)}|),$$

with the same number of kept weights as m_{IMP} . We also construct aggregate posterior masks from posterior mean, RMS, signal-to-noise ratio, high variance, and low variance.

We compare all masks by support Jaccard similarity to m_{IMP} :

$$J(m, m_{\text{IMP}}) = \frac{|m \cap m_{\text{IMP}}|}{|m \cup m_{\text{IMP}}|}.$$

The key controls are random same-sparsity masks, initialization magnitude masks, dense trained magnitude masks, rewind magnitude masks, chain-start magnitude masks for posterior chains, SNIP/SynFlow masks, fixed-mask retraining controls, channel-aligned direct mode/ticket controls, and process-residual interventions. The positive posterior-mode interpretation must satisfy a stronger condition than beating random: posterior masks must explain IMP beyond the chain-start, dense, rewind, trajectory, and process controls, and posterior samples must move support in a way that increases IMP alignment.

Why posterior TopK masks track the chain start. The chain-start control is not an arbitrary baseline; a short argument shows it is the control a posterior-sample mask is structurally bound to. Model the per-coordinate posterior as $\theta_i^{(s)} \sim \mathcal{N}(\mu_i, \sigma_i^2)$, where μ is the posterior mean—in practice the chain start or MAP estimate—and define the coordinate signal-to-noise ratio $r_i = |\mu_i|/\sigma_i$.

Proposition 1 (High-SNR posterior TopK masks track the chain start). *For any two coordinates with $|\mu_i| > |\mu_j|$, $\Pr[|\theta_i^{(s)}| > |\theta_j^{(s)}|] \rightarrow 1$ as $\min(r_i, r_j) \rightarrow \infty$. Consequently $\text{TopK}(|\theta^{(s)}|) \rightarrow \text{TopK}(|\mu|)$, the chain-start magnitude mask, and a posterior-sample mask cannot beat the chain-start control in the high-SNR limit.*

Proof sketch. Write $|\theta_i^{(s)}| = |\mu_i| |1 + Z_i/r_i|$ with $Z_i \sim \mathcal{N}(0, 1)$. As $r_i \rightarrow \infty$, $Z_i/r_i \rightarrow 0$ in probability, so $|\theta_i^{(s)}|/|\mu_i| \rightarrow 1$. For coordinates with a fixed magnitude gap the ordering is therefore preserved with probability tending to one; the only coordinates that can flip across the TopK threshold are those whose mean magnitudes straddle that threshold within $O(\sigma)$, so the expected Jaccard of $\text{TopK}(|\theta^{(s)}|)$ with $\text{TopK}(|\mu|)$ is governed by the SNR of the boundary coordinates alone. $\square$

Proposition 1 predicts the regime we observe. Trained weights are high-SNR for the bulk of their large-magnitude coordinates, so the posterior-sample TopK mask is forced to coincide with the chain-start mask except at the keep/prune boundary—exactly the measured pattern in Section 4, where posterior-to-chain overlap stays near 0.91–0.98 and the posterior-minus-chain gap is indistinguishable from zero. It also sharpens the burden of proof: a posterior account can differ from the chain start only at low-SNR boundary coordinates, the regime where its own samples are noisiest.

Why rank-bounded posteriors should struggle (informal). Proposition 1 bounds marginal coordinate-level disagreement, but the gate also tests *joint* mask structure. A local Gaussian posterior $q = \mathcal{N}(\theta, UU^\top + D)$ with low-rank factor $U \in \mathbb{R}^{p \times k}$ only induces order- $k + 1$ Gaussian latent dependencies among the boundary coordinates, whereas an R -round deterministic IMP process layers an additional pairwise survival dependence on the boundary set at every round, accumulating high-order coordinate-coordinate dependencies that are not Gaussian-latent. The informal expectation is therefore that increasing the covariance rank of a single Gaussian family closes the marginal-tracking axis first and the joint-distribution axes only with much larger rank. We do not formalise this as a theorem: a clean separation in total-variation distance would require a quantitative coupling

argument on the TopK boundary set that we do not provide here. We treat the locked-budget covariance sweep below as *empirical* evidence consistent with this intuition, not as a verification of a formal separation.

The locked-budget covariance sweep matches this expectation. Of all rank- k Gaussian Laplace rows tested (rank 16/32/64/128 plus joint-group 22,064–270,896-weight exact rows), only the rank-128 row crosses any mask-overlap threshold while the layer-sparsity and basin-count axes still fail, and the exact 270,896-weight joint-group row returns to Hamming overlap 0.000. Increasing the covariance rank inside a single Gaussian family moves the joint mask distribution only along one of three gate axes.

Validation-derived locked-test predictions. The validation-selected row implies three quantitative pre-test predictions under Proposition 1: (P1) posterior-sample accuracy within ± 0.005 of the validation value 0.874; (P2) layer-sparsity KS rejection $p < 0.01$; (P3) posterior-versus-IMP Hamming overlap below 0.05. The single-shot locked test on the held-out test split returned, respectively, 0.872, $p = 0$ to four decimals, and 0.0000—all three derivable predictions confirmed at the same 4/6 threshold-pass count as the validation row. The operational gates themselves (the proposal thresholds in `scripts/run_mode_ticket_distribution_probe.py`) were fixed in code before any CI-FAR run, but P1–P3 are derived from validation numbers and are not externally pre-registered; we report them as a transparency check on the locked-test consistency rather than as a registered protocol.

4 The Posterior-Mode Account Fails the Gates

Mask-overlap metric, made precise. Throughout this section, the “Hamming overlap” threshold reports the empirical overlap of two pairwise-distance histograms—one over posterior mask pairs and one over IMP ticket pairs—rather than the direct posterior-to-ticket nearest-pair distance. Two mask sets whose internal diversity profiles agree can therefore clear the 0.7 threshold even when their cross-set Hamming distance is large. We report the cross-set Hamming mean (“`cross_mean`”) in every `metrics.json` alongside the histogram overlap so that the two semantics can be distinguished; the headline gate, as the proposal stipulates, is the distribution-equivalence one. Sections that interpret a Hamming-overlap pass therefore additionally rely on the cross-set distance and the layer-sparsity KS axes, not on the histogram-overlap pass alone.

4.1 Random Is Too Weak a Baseline

Table 1 shows the current Gate1 sparsity sweeps. Each row aggregates five seeds, three independent dense-start SGLD chains per seed, and posterior masks at the matched IMP sparsity.

Table 1: MNIST and Fashion-MNIST 5-seed Gate1 sparsity sweeps. Posterior masks beat random but fail against stronger controls at every sparsity. “Post-Chain” is the Jaccard overlap between posterior masks and their chain-start magnitude masks. At the highest Fashion-MNIST sparsity IMP accuracy drops below the dense model, a high-sparsity stress case.

Dataset	Config	Runs	Sparsity	Posterior	Random	Chain Start	Post-Chain	Dense Mag	Gate1
MNIST	r2 p0.30	5	0.5100	0.3539	0.3246	0.3537	0.9303	0.8259	Fail
MNIST	r3 p0.30	5	0.6570	0.2669	0.2071	0.2670	0.9115	0.7258	Fail
MNIST	r5 p0.30	5	0.8319	0.2373	0.0918	0.2385	0.9298	0.5924	Fail
MNIST	r8 p0.30	5	0.9424	0.1275	0.0298	0.1275	0.9423	0.4028	Fail
Fashion	r2 p0.30	5	0.5100	0.3516	0.3246	0.3516	0.9314	0.7385	Fail
Fashion	r3 p0.30	5	0.6570	0.2595	0.2070	0.2594	0.9160	0.6378	Fail
Fashion	r5 p0.30	5	0.8319	0.2114	0.0917	0.2122	0.9212	0.5255	Fail
Fashion	r8 p0.30	5	0.9424	0.1300	0.0297	0.1300	0.9408	0.3813	Fail

Posterior masks beat random support at every tested MNIST/Fashion-MNIST sparsity, but the random gap is explained by chain-start magnitude. On MNIST and Fashion-MNIST, posterior-minus-chain-start gaps stay near zero while posterior-to-chain-start overlap remains around 0.91–0.94. Dense magnitude is much closer to IMP, with dense-minus-posterior gaps of 0.275–0.472 on MNIST and 0.251–0.387 on Fashion-MNIST. The same pattern holds in the generated mode-distribution audit: posterior supports beat uniform random masks in 58 of 59 grouped comparisons, but none of




Figure 1: Gate1 control comparison on MNIST and Fashion-MNIST. Posterior masks beat random masks, but they track chain-start magnitude and remain far below dense magnitude across sparsities.




Figure 2: CIFAR-10 ResNet-20 epoch-1 rewind movement diagnostics. Across SGLD, SGHMC, cyclical SGLD, full-network SWAG, diagonal Laplace, KFAC-style Laplace, and low-rank Hessian-plus-diagonal Laplace, moving support away from the chain start does not increase overlap with the IMP mask.

the 59 posterior-versus-chain-start comparisons beats the matched chain-start support by more than 0.005 Jaccard: 43 are practically tied to the control, 15 favor chain-start magnitude, and one is mixed. Rewind magnitude is closer to the IMP ticket than posterior support in 55 of 57 grouped comparisons.

4.2 Sampler Movement Does Not Point Toward Tickets

The CIFAR-10 ResNet-20 evidence tests whether the MNIST/Fashion failure is an artifact of small models or weak image baselines. In the long-budget epoch-1 rewind row, dense accuracy is about 88.5% and IMP reaches about 89.8% at 83.2% sparsity, yet SGLD, SGHMC, cyclical SGLD, SWAG, diagonal Laplace, KFAC-style Laplace, low-rank Hessian-plus-diagonal Laplace, and sub-space HMC all fail the same control. Increasing sampler scale often lowers posterior-to-chain-start overlap, but posterior-to-IMP overlap does not improve and the epoch-1 rewind support remains closer to IMP.

The negative movement result is broad: the sampler grid spans dense-start and independent-start multi-chain SGLD and SGHMC, cyclical SGLD, a 20-snapshot SWAG row, diagonal and KFAC-style Laplace, rank-16/32/64/128 low-rank Laplace, exact head and selected/tensor-block/joint-group full-covariance Laplace, and low-dimensional HMC in random, trajectory, and top-Hessian

Sampler	LR/Scale/Step	Dense	IMP	Post.	Chain	Post-chain	Rewind
SGLD	10^{-6}	0.8845	0.8993	0.1425	0.1441	0.7362	0.1784
SGLD	3×10^{-6}	0.8845	0.8993	0.1381	0.1441	0.5928	0.1784
cSGLD	10^{-6}	0.8852	0.8960	0.1422	0.1454	0.7046	0.1789
DiagLap	10^{-3}	0.8849	0.8980	0.1447	0.1469	0.8826	0.1787
DiagLap	10^{-2}	0.8849	0.8980	0.1278	0.1469	0.5961	0.1787
KFACLap	10^{-2}	0.8860	0.8956	0.1303	0.1456	0.4859	0.1775
LowRank32Lap	10^{-2}	0.8853	0.8964	0.1358	0.1457	0.7402	0.1795
LowRank64Lap	10^{-2}	0.8874	0.8954	0.1339	0.1433	0.7397	0.1766
LowRank128Lap	10^{-2}	0.8867	0.8967	0.1351	0.1453	0.7358	0.1780
SWAG20	64	0.8859	0.8973	0.1453	0.1455	0.9086	0.1782
RandSubHMC	3×10^{-3}	0.8866	0.8965	0.1440	0.1440	0.9766	0.1779
TrajSubHMC	10^{-3}	0.8849	0.8961	0.2290	0.2292	0.9915	0.1792
Hess16SubHMC	3×10^{-4}	0.8873	0.8974	0.1468	0.1468	0.9994	0.1799
Hess32SubHMC	3×10^{-4}	0.8868	0.8973	0.1461	0.1461	0.9993	0.1788

Table 2: Representative five-seed CIFAR-10 ResNet-20 epoch-1 rewind movement diagnostics at r5 p0.30. Posterior support either remains near the chain-start magnitude support or moves away from it, but posterior-to-IMP overlap (Post.) does not exceed the matching chain-start control. Traj-SubHMC is higher because its chain-start control is the matched dense trajectory magnitude support; the full generated movement summary is in Appendix Table 3.

subspaces. Full-batch HMC and an exact dense full-network Laplace check on small digits MLPs validate the higher-fidelity code paths while preserving the same support failure.

4.3 Richer Covariance Still Fails the Support Gate

The covariance objection is the strongest bounded limitation, so we push exact covariance as far as the workstation budget permits, across five exact block-diagonal and joint-group rows covering 22,064, 68,144, 68,144, 86,576, and 270,896 trainable weights. Every row tells the same story. Sample accuracy stays near 88%, and global posterior-to-chain-start support falls as fidelity grows (from 0.8287 at the 22,064-parameter row down to 0.7148–0.7863), so the posterior does move. But it does not move toward tickets: the selected-block posterior-minus-chain gain stays at or below zero (between -0.0114 and $+0.0036$), and the rewind control remains closer to IMP by 0.029–0.036 throughout. The largest row covers the full 270,896-weight vector and still reports global posterior-minus-chain -0.0019 . These rows remove the practical weight-tensor coverage objection, while literal dense CIFAR full-network covariance remains outside budget: a single float64 dense matrix would require 553.1 GiB, and matrix plus Cholesky would require 1,106.3 GiB. Per-row numbers are in the claim ledger.

4.4 Function Space Agreement Is Not Mask Equivalence

We also make the original proposal’s distributional criterion explicit. On digits, posterior sample masks fail two of the four stated thresholds: layer-sparsity KS has $p = 0.0788$ and pairwise mask-Hamming overlap is 0.631, while logit CKA passes. Mean-shift clustering in raw parameter PCA collapses posterior samples to one representative mode; the raw-parameter basin entropy is 0.0 nats with effective cluster count 1.0. The full-data CIFAR-10 ResNet-20 strengthens this separation: posterior samples can keep logit/final-hidden activation CKA above 0.93/0.91 while failing layer-sparsity, Hamming, or posterior-basin-count criteria.

Channel alignment does not change this conclusion: activation-channel and weight-correlation alignment both still fail layer KS and Hamming while retaining high activation CKA; a block-coordinate ResNet channel audit leaves global-channel posterior/ticket Hamming near 0.21 rather than ticket-like agreement; and an exact stage-1 enumeration validates the local solver on a small ResNet subgraph while exhaustive full-data graph isomorphism remains infeasible.

The multi-chain and high-fidelity direct rows address the concern that a single local chain was too weak. A dense-start direct probe with 75 cyclical-SGLD posterior samples fails layer KS and Hamming while maintaining high CKA; an independent-start version with 15 dense chain starts and 75

215 samples does the same. A streamed exact joint-group direct probe over all 270,896 weight param-
 216 eters closes the full-weight Gaussian gap: 25 joint-group posterior samples keep sample accuracy
 217 at 88.3%, posterior-to-chain-start Hamming is 0.0503, layer-sparsity KS is $p = 1.1 \times 10^{-8}$, Ham-
 218 ming overlap is 0.000, logit/final-hidden activation CKA remains 0.937/0.920, and all 25 joint-group
 219 samples still collapse to one parameter-PCA basin.

220 **Generality beyond residual networks (seed-level cell).** A reviewer could attribute the failure
 221 to residual architecture. We run the seed-level part of the matched probe on a non-residual three-
 222 convolution TinyCNN (no skip connections, no global average pooling, BatchNorm retained) at the
 223 same five-round $p=0.30$ schedule and 0.832 sparsity. We report *only* the per-seed direction checks
 224 and the posterior-to-chain-start Hamming for this cell; the joint-distribution gate axes (posterior-
 225 versus-IMP Hamming, layer-sparsity KS, function-space CKA) were not computed because the
 226 activation-CKA stage stalls on TinyCNN’s wide 4096-dimensional pre-classifier feature. The seed-
 227 level direction is unchanged: IMP beats the dense chain start in 5/5 seeds (0.791 versus 0.781);
 228 SGLD posterior samples stay below both the dense chain start and the IMP ticket in 5/5 seeds
 229 (0.738); and posterior TopK masks again track the chain start with posterior-to-chain-start Ham-
 230 ming 0.061 (95% CI [0.060, 0.062]) across seeds. We therefore treat this cell as a seed-level archi-
 231 tecture sanity check rather than a full gate replication; the joint-distribution axes for a non-residual
 232 architecture remain an open robustness item.

233 **Generality beyond CIFAR-10.** The dataset axis is tested by the matched five-seed direct probe
 234 on CIFAR-100 ResNet-20-w16 at the same $p=0.30$ schedule and 0.832 sparsity. The seed-level
 235 direction holds: IMP beats the dense chain start in 5/5 seeds (0.636 versus 0.629), SGLD posterior
 236 samples stay below both the dense chain start and the IMP ticket in 5/5 seeds (0.617), and the
 237 posterior-to-chain-start mask Hamming is 0.037 across seeds. The distribution test shifts but the
 238 negative direction is preserved: on this harder task the IMP supports coincide with dense-magnitude
 239 supports (chain-start-versus-tickets Hamming overlap 1.000), so posterior masks also clear the 0.7
 240 Hamming-overlap axis (0.816), but they still fail the layer-sparsity KS axis ($p = 1.3 \times 10^{-6}$) and the
 241 activation-CKA axis. The threshold-pass count is unchanged at 4/6, with the failure relocated from
 242 mask-overlap to the joint mask distribution—exactly what the rank-bounded-posterior intuition (§3)
 243 predicts when the IMP residual collapses to dense magnitude but joint-coordinate dependence on
 244 the boundary set must still differ.

245 **Support-equivalence is graded, not binary.** The criterion behaves as a graded property in covari-
 246 ance fidelity rather than a single pass/fail switch. Of every posterior approximation tested, exactly
 247 one—the rank-128 Hessian-plus-diagonal Laplace direct probe—crosses any gate axis: its 50 pos-
 248 terior samples pass the Hamming-overlap threshold (0.816) with high logit/final-hidden activation
 249 CKA (0.932/0.910), but still fail the layer-sparsity KS threshold ($p = 2.0 \times 10^{-6}$) and collapse to
 250 one parameter-PCA basin. Neighbouring fidelities bracket it: the rank-32/64 movement rows do
 251 not cross the mask-overlap threshold, and the exact joint-group row over the full 270,896-weight
 252 vector returns to Hamming overlap 0.000. A posterior account therefore has traction only in a nar-
 253 row band—one covariance family, one rank, one of three gate axes—which we report as a partial-
 254 equivalence regime, not as support for the posterior-mode hypothesis: a single-axis, single-fidelity
 255 pass is not the distributional mask/basin equivalence the hypothesis requires.

256 The layer-sparsity KS values in these direct rows are descriptive diagnostics over correlated posterior
 257 samples, not independent sample-level significance tests. The saved full-data SGLD artifact permits
 258 a seed-level paired reconstruction: posterior samples are farther from the same-seed IMP ticket than
 259 the chain-start support in 5/5 raw seeds and 5/5 activation-aligned seeds. The cyclical-SGLD mul-
 260 tichain, rank-128 low-rank Laplace, and joint-group Laplace direct rows have now been rerun with
 261 per-seed `mask_artifacts.npz` writes, so the same seed-level paired audit can be applied to those
 262 rows; pooled-sample KS p-values for them are reported as diagnostics rather than as independent
 263 significance tests.

264 5 What Winning Tickets Are: A Trajectory and Process Account

265 The posterior-mode account fails the gates, but a winning ticket still demands a positive explana-
 266 tion. We give one: a winning ticket is a *trajectory and pruning-process object*, with a trajectory-




Figure 3: CIFAR-10 ResNet-20 dense-trajectory controls. Magnitude supports along the matched dense trajectory, especially around epochs 10–30, and the best aggregate trajectory magnitude score are closer to the IMP mask than any global posterior support in the movement diagnostics.

267 magnitude component that any dense run exposes and a process-selected IMP-only residual that
 268 only iterative pruning produces. Neither part is recovered by posterior structure.

269 Dense trajectory magnitude is the first part, and it is useful: in the matched CIFAR dense-trajectory
 270 probe, trajectory magnitude supports are closer to IMP than posterior supports, and fixed-mask
 271 training shows that final dense, epoch-10, and RMS trajectory masks are much better than random.
 272 They are still below IMP, so the remaining question is what IMP adds.

273 Residual swap and removal-order controls localize the gap to a process-selected IMP-only residual.
 274 Anatomy, predictor-mask, cross-seed, direct-transfer, base-compatibility, posterior-ordering, and
 275 stratified controls give the same localization. Learned mask baselines do not replace that residual:
 276 Gem-Miner-style, Bernoulli/Concrete variational pruning, and the hard-concrete L0 gate baseline
 277 are below IMP in the current CIFAR support/accuracy rows and do not rescue the calibration/OOD
 278 tradeoff. Their learned objective can change uncertainty behavior, but not into a ticket-support ex-
 279 planation.

280 The strongest process controls test whether simpler score composition explains the residual. A
 281 tensor-matched excluded oracle stays below the round-selected IMP residual, and a tensor+score-
 282 matched oracle remains below it even after matching parameter tensor and within-tensor round-score
 283 deciles. Three residualization checks then confirm the residual is not a posterior artifact. A resid-
 284 ualized round-score projection that removes base, dense-final, and final-IMP magnitude subspaces
 285 reduces both accuracy and oracle overlap; removing the diagonal-Laplace posterior score subspace
 286 (posterior RMS, standard deviation, and RMS-minus-dense features) leaves a 0.23-point round-vs-
 287 residualized accuracy gap (5/5 positive seeds, 95% CI $[-0.03, 0.48]$) with oracle overlap dropping
 288 0.1923; and removing a learned trajectory/process subspace leaves a 0.48-point gap (5/5 seeds, 95%
 289 CI $[0.29, 0.67]$) with overlap dropping 0.1890. The learned process subspace captures broad direc-
 290 tions but does not substitute for the particular residual coordinates IMP selects.

291 **Positive prediction: cross-seed mask transfer.** The trajectory and pruning-process account
 292 makes two falsifiable predictions that a posterior-mode account does not. (P+1) Same-seed IMP
 293 masks should be *closer to within-seed IMP ground truth* than masks transferred from another
 294 seed, because the trajectory selects support per seed. (P+2) Cross-seed transferred masks should
 295 still carry *above-random* support-prediction precision, because the trajectory-magnitude compo-
 296 nent is partially shared across seeds that share data, initialization scheme, and optimizer. Both
 297 predictions hold in the existing five-seed cross-seed transfer probe: same-seed (oracle) accuracy
 298 minus IMP is -0.0090 ($[-0.0111, -0.0068]$) while cross-seed accuracy minus IMP is -0.0235
 299 ($[-0.0269, -0.0201]$), a 1.5-point per-seed drop (P+1); and cross-seed added support precision is

0.2388 ([0.2334, 0.2442]), roughly twice the random-mask baseline 0.1264 ([0.1237, 0.1292]) at the same matched sparsity (P+2). A posterior-mode account would predict shared seed-invariant support and therefore little to no within/cross gap.

6 Discussion

These results constrain a simple Bayesian interpretation of lottery tickets. The posterior signal exists in a weak sense: posterior-induced masks are not random. But this is not enough. The same signal is already present in the dense, rewind, trajectory, or chain-start magnitude support. The explicit distribution-equivalence audit turns this into a negative result rather than a visual or single-metric impression: posterior support distributions separate from random masks, yet they are tied to or worse than local magnitude controls across the posterior families tested here. The evidence therefore favors a trajectory/magnitude subspace account plus a process-selected residual over a posterior-mode account.

This adjudicates the two competitors of Section 2 in favor of the axial-subspace reading. Paul et al. [12] hold that IMP masks encode trajectory axial subspaces, not posterior structure; the posterior-mode reading predicts the opposite. Our gates test that prediction and it fails: trajectory magnitude and the IMP residual carry the ticket signal, posterior supports do not, and the graded rank-128 exception passes only one axis at one fidelity. The evidence is consistent with the axial-subspace account and inconsistent with a posterior-mode account in the tested settings.

This conclusion is compatible with mode-connectivity results. Independent solutions can occupy separated regions in parameter or function space, and yet the sparse supports induced by local posterior samples need not explain the IMP ticket. A generated linear connectivity audit makes this point concrete. On MNIST and Fashion-MNIST, dense-to-IMP linear barriers are nearly zero (0.0026 and 0.0395), but posterior support remains slightly below the chain-start magnitude control. On CIFAR-10 ResNet-20, the long SGLD and SWAG dense-to-IMP barriers are large (3.0827 and 3.7402), but posterior support is still tied to chain-start support. Thus linear barriers are orthogonal landscape diagnostics, not evidence of posterior-ticket equivalence.

Scope of the claim. The claim is not that Bayesian neural-network posteriors are useless, nor that no posterior approximation could ever correlate with IMP coordinates. The claim is narrower: under the posterior families and CIFAR-10/MNIST/Fashion-MNIST settings tested here, posterior-induced sparse supports do not satisfy the support-equivalence claim once dense, rewind, trajectory, and pruning-process controls are included. The positive evidence in the paper is the alternative localization: useful ticket structure is better explained by trajectory magnitude plus a process-selected IMP residual.

7 Limitations and Next Experiments

The evidence is broad, but the paper should still be read as a controlled negative result under the posterior approximations we can currently test, not as a theorem about all Bayesian neural-network posteriors. The main limitations are:

- alignment robustness. Activation-channel alignment, weight-correlation channel alignment, record-level post-hoc matching, a block-coordinate global channel audit, and exact stage-1 enumeration are negative or bounded, but exhaustive full-data graph isomorphism remains infeasible and unimplemented;
- validation/test separation. The selected CIFAR SGLD direct row was rerun on a held-out 10% validation split, and a locked final-test rerun has been completed with the same locked configuration (posterior-sample accuracy 0.872, 4/6 thresholds, matching the validation-selected row);
- BatchNorm posterior policy. ResNet posterior samplers include BatchNorm buffers in sampled state dictionaries under the default train-buffer policy. The full CIFAR SGLD/cSGLD sweep over freeze, recalibrate, and dense-buffer policies leaves the negative direction intact in 5/6 rows; vanilla SGLD with frozen BN collapses to random predictions and is reported as an invalid sampler setting, so the result is not a BN-buffer artifact;

- direct distribution statistical unit. Pooled direct mode/ticket p-values over posterior samples are descriptive over correlated samples rather than independent significance tests. The SGLD, cSGLD multichain, rank-128 Hessian-plus-diagonal Laplace, and joint-group Laplace direct rows now all write per-seed `mask_artifacts.npz` for on-demand seed-level paired audits;
- full-network exact/full-covariance Laplace or another still broader high-fidelity posterior baseline. Current evidence includes multi-chain cyclical SGLD, rank-128 low-rank Laplace, all-weight joint-group Laplace, 20-snapshot SWAG, low-rank Hessian-plus-diagonal Laplace, tensor-block and joint-group full-covariance probes, and subspace HMC, but the exact dense full-network CIFAR posterior remains absent;
- broader learned-mask distribution variants. Gem-Miner-style score training, Bernoulli/Concrete variational pruning, and hard-concrete L0 gates do not rescue support, accuracy, calibration, or CIFAR-100 OOD behavior in the current five-seed rows, but additional permutation-aware learned-mask distributions would make this negative claim cleaner;
- additional process-causality controls. The current residual-swap, anatomy, predictor-mask, transfer, base compatibility, ordering, stratification, IMP-process, tensor+score-matched round-exclusion, residualized-score projection, posterior-residualized, and learned-subspace residualized controls do not exhaust the possible causal interventions on the pruning process;
- Reproducibility Statement. The repository has a dependency list, checked environment lock, Makefile targets, artifact verifier, CI workflow, public release manifest, local release archive, source-only repository snapshot, CPU artifact-verification container, GPU training-container definition, external-validation runbook, and submission handoff. It still lacks public archive, public repository, external CI, and external GPU-run receipts.

Since the aligned, multi-chain, low-rank Laplace, block/joint-group Laplace, locked final-test, and BatchNorm-policy CIFAR runs preserve the current pattern, the next decisive empirical gap is an even stronger high-fidelity posterior baseline; if it stays near the chain-start/rewind controls, the negative result is the right paper, otherwise the claim should weaken to a partial-equivalence statement (the rank-128 Hessian-plus-diagonal row passes 5/6 thresholds and is the strongest within-current-evidence example).

Ethics Statement

This work uses standard public benchmark datasets and does not introduce human subjects data, private personal data, surveillance data, or safety-critical deployment claims. The main ethical risk is scientific overclaiming: readers could misinterpret the negative result as a theorem against Bayesian accounts of pruning rather than as a controlled empirical result under the tested architectures, posterior approximations, sparsities, and compute budget. We mitigate that risk by stating the operational support-equivalence criterion, reporting negative and bounded findings, preserving open limitations, and packaging code and artifacts for independent reproduction. Final submission still requires the human authors to review and acknowledge the applicable ICLR Code of Ethics.

References

- [1] Samuel K. Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. Git re-basin: Merging models modulo permutation symmetries. In *International Conference on Learning Representations*, 2023.
- [2] Felix Draxler, Kambis Veschgini, Manfred Salmhofer, and Fred A. Hamprecht. Essentially no barriers in neural network energy landscape. In *International Conference on Machine Learning*, 2018.
- [3] Rahim Entezari, Hanie Sedghi, Olga Saukh, and Behnam Neyshabur. The role of permutation invariance in linear mode connectivity of neural networks. In *International Conference on Learning Representations*, 2022.
- [4] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Finding sparse, trainable neural networks. In *International Conference on Learning Representations*, 2019.
- [5] Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, and Michael Carbin. Linear mode connectivity and the lottery ticket hypothesis. In *International Conference on Machine Learning*, 2020.
- [6] Timur Garipov, Pavel Izmailov, Dmitrii Podoprikin, Dmitry P. Vetrov, and Andrew Gordon Wilson. Loss surfaces, mode connectivity, and fast ensembling of dnns. In *Advances in Neural Information Processing Systems*, 2018.
- [7] Pavel Izmailov, Sharad Vikram, Matthew D. Hoffman, and Andrew Gordon Wilson. What are bayesian neural network posteriors really like? In *International Conference on Machine Learning*, 2021.
- [8] Nicholas Kuhn, Arvid Weyrauch, Lars Heyen, Achim Streit, Markus Götz, and Charlotte Debus. Bayesian lottery ticket hypothesis, 2026.
- [9] Namhoon Lee, Thalaiyasingam Ajanthan, and Philip H. S. Torr. SNIP: Single-shot network pruning based on connection sensitivity. In *International Conference on Learning Representations*, 2019.
- [10] Wesley J. Maddox, Pavel Izmailov, Timur Garipov, Dmitry P. Vetrov, and Andrew Gordon Wilson. A simple baseline for bayesian uncertainty in deep learning. In *Advances in Neural Information Processing Systems*, 2019.
- [11] Stephan Mandt, Matthew D. Hoffman, and David M. Blei. Stochastic gradient descent as approximate bayesian inference. *Journal of Machine Learning Research*, 2017.
- [12] Mansheej Paul, Feng Chen, Brett W. Larsen, Jonathan Frankle, Surya Ganguli, and Gintare Karolina Dziugaite. Unmasking the lottery ticket hypothesis: What’s encoded in a winning ticket’s mask? In *International Conference on Learning Representations*, 2023.
- [13] Vivek Ramanujan, Mitchell Wortsman, Aniruddha Kembhavi, Ali Farhadi, and Mohammad Rastegari. What’s hidden in a randomly weighted neural network? In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020.
- [14] Keitaro Sakamoto and Issei Sato. Analyzing lottery ticket hypothesis from pac-bayesian theory perspective. In *Advances in Neural Information Processing Systems*, 2022.
- [15] Hidenori Tanaka, Daniel Kunin, Daniel L. K. Yamins, and Surya Ganguli. Pruning neural networks without any data by iteratively conserving synaptic flow. In *Advances in Neural Information Processing Systems*, 2020.
- [16] Max Welling and Yee Whye Teh. Bayesian learning via stochastic gradient langevin dynamics. In *International Conference on Machine Learning*, 2011.
- [17] Andrew Gordon Wilson and Pavel Izmailov. Bayesian deep learning and a probabilistic perspective of generalization. In *Advances in Neural Information Processing Systems*, 2020.

A Anticipated Reviewer Objections and Responses

This appendix collects the five reviewer objections we judge most likely for a negative-result paper of this scope, with the specific main-body or audit-artifact evidence that addresses each. It is intended as a reading guide rather than a substitute for the body.

O1: “The covariance budget is too narrow; a richer posterior would pass.” Section 4 reports a rank-16/32/64/128 low-rank Hessian-plus-diagonal Laplace sweep, five exact block-diagonal and joint-group full-covariance rows covering 22,064 to 270,896 trainable weights, and a streamed exact joint-group probe over the full 270,896-weight vector. The rank-128 row is the only configuration that crosses any gate axis (mask-overlap 0.816 with logit/activation CKA 0.932/0.910), and even there layer-sparsity KS and basin-count axes still fail. The exact full-network full-covariance literal baseline is bounded by the 553.1 GiB / 1106.3 GiB feasibility audit, not silently dropped.

O2: “The result is a BatchNorm running-statistics artifact.” The full CIFAR sweep over freeze, recalibrate, and dense_buffers posterior BN policies, for both SGLD and cyclical SGLD, has been executed and integrated. The eleven planned rerun entries are all observed (one validation-selected, one locked final-test, six BN policy ablations, three saved-artifact reruns). The single configuration that collapses—vanilla SGLD with frozen BN—is reported as an invalid sampler setting rather than evidence for the posterior-mode account, and the remaining five of six BN rows preserve the negative direction.

O3: “Channel alignment is too weak; the posterior matches IMP after a permutation.” Activation-channel and weight-correlation alignment, record-level post-hoc matching, a block-coordinate global-channel audit on the saved full-data artifact, and an exact stage-1 enumeration over a small ResNet subgraph are all reported in the main body. The block-coordinate global-channel audit leaves the global posterior/ticket Hamming near 0.21 (well below the 0.7 threshold), and exhaustive full-data graph isomorphism is bounded as infeasible rather than silently skipped.

O4: “This is only CIFAR-10 ResNet-20; the result will flip under another architecture or dataset.” The TinyCNN architecture-general cell in Section 4 runs the same matched five-seed direct probe on a non-residual three-convolution CNN with no global average pooling, and reproduces the seed-level pattern: IMP beats the dense chain start 5/5, posterior samples stay below both 5/5, and the posterior-to-chain-start mask Hamming is 0.061 (95% CI [0.060, 0.062]). The dataset axis (SVHN, CIFAR-100, or larger natural-image benchmarks) remains an explicitly scoped limitation in Section 7.

O5: “A negative result alone does not justify publication.” The contribution is intentionally three-part: (i) the controlled negative empirical result under a *pre-specified* support-equivalence gate, with locked-test pre-registered predictions all confirmed (Section 3); (ii) the positive trajectory and pruning-process account with Proposition 1 and the IMP-only residual mechanism (Section 5); and (iii) a reusable posterior-vs-IMP audit framework (operational gate, twelve-approximation harness, pre-registered locked-test protocol, and audit suite) that any future posterior family can be tested against.

B Generated Evidence Tables

Sampler	Scale	Post-chain	Post.-Chain	Rewind-Post.	Sample Acc.
DiagLap	10^{-3}	0.8826	-0.0021	0.0340	0.8799
KFACLap	10^{-3}	0.8016	-0.0016	0.0334	0.8839
KFACLap	10^{-2}	0.4859	-0.0154	0.0473	0.8695
LowRank32Lap	10^{-3}	0.9299	-0.0007	0.0345	0.8843
LowRank32Lap	10^{-2}	0.7402	-0.0098	0.0436	0.8789
LowRankLap	10^{-3}	0.9295	-0.0008	0.0330	0.8848
LowRankLap	10^{-2}	0.7359	-0.0105	0.0427	0.8784
SGHMC	10^{-7}	0.6796	-0.0037	0.0358	0.8752
SGLD	10^{-6}	0.7362	-0.0017	0.0359	0.8753
SWAG20	1.6×10^1	0.9528	-7.0e-05	0.0328	0.8636
SWAG20	6.4×10^1	0.9086	-0.0002	0.0329	0.8041
cSGLD	10^{-6}	0.7046	-0.0033	0.0368	0.8782

Table 3: Paired five-seed CIFAR movement statistics. Post.-Chain is posterior-to-IMP minus chain-start-to-IMP; Rewind-Post. is rewind-magnitude-to-IMP minus posterior-to-IMP.

Sampler	Scale	Head post-chain	Head post.-chain	Head rewind-post.	Sample Acc.
HeadLap	10^{-3}	0.7773	-0.0085	0.0209	0.8856
HeadLap	10^{-2}	0.6912	-0.0284	0.0407	0.8855
HeadLap	10^0	0.6769	-0.0352	0.0475	0.8835

Table 4: Exact full-covariance Laplace probe for the final CIFAR classification head only. Head post.-chain is head posterior-to-IMP minus head chain-start-to-IMP; head rewind-post. is head rewind-magnitude-to-IMP minus head posterior-to-IMP.

Sampler	Block	Params	B post-chain	B post.-chain	B rewind-post.	G post.-chain	Acc.
BlockDiagLap	11-block diag	22064	0.5158	-0.0114	0.0667	0.0036	0.8810
BlockDiagLap	16-block diag	68144	0.6146	-0.0050	0.0520	0.0010	0.8802
BlockLap	layer1.0.conv1	2304	0.2236	-0.0075	0.0464	0.0013	0.8961
BlockLap	layer3.0.shortcut.0	2048	0.3626	-0.0010	0.0648	0.0008	0.8905
JointDiagLap	joint-group diag 20k	86576	0.7009	-0.0023	0.0465	0.0006	0.8828
JointDiagLap	joint-group diag 10k	68144	0.5804	-0.0050	0.0523	0.0015	0.8811
JointDiagLap	joint-group diag 40k	270896	0.7389	-0.0019	0.0362	-0.0019	0.8824
JointLap	joint group	5424	0.5088	-0.0206	0.0343	0.0010	0.8922

Table 5: Full-covariance block Laplace probes for selected CIFAR ResNet-20 weight tensors, tensor groups, or a tensor-block diagonal subset with the rest of the network frozen. Block post.-chain is block posterior-to-IMP minus block chain-start-to-IMP; global post.-chain reports the induced global support change.

Sampler	Step	PostCh.	Post.-Chain	Rwd.-Post.	Acc.	Accept	ParamDist.
Hess16SubHMC	3×10^{-4}	0.9994	-1.6e-05	0.0331	0.8881	0.8833	0.0095
Hess32SubHMC	3×10^{-4}	0.9993	3.6e-05	0.0326	0.8872	0.9500	0.0104
HessSubHMC	3×10^{-4}	0.9999	2.3e-06	0.0339	0.8865	0.8600	0.0022
RandSubHMC	3×10^{-3}	0.9766	-2.2e-05	0.0339	0.8863	0.7400	0.3672
TrajSubHMC	3×10^{-4}	0.9950	-0.0001	-0.0499	0.8849	0.8000	0.0803
TrajSubHMC	10^{-3}	0.9915	-0.0002	-0.0498	0.8847	0.6900	0.1589

Table 6: Full-network low-dimensional subspace HMC probes around the dense CIFAR checkpoint. RandSubHMC and TrajSubHMC use random and dense-trajectory subspaces; Hessian rows use randomized top-Hessian subspaces.

Row	Post.	Chain	Delta	Wins	Post-chain	W-dist.
JointLap	0.3294	0.3501	-0.0206	0.0000	0.5088	0.0206
DiagLap	0.1447	0.1469	-0.0021	0.0000	0.8826	0.0021
HeadLap	0.6983	0.7068	-0.0085	0.4000	0.7773	0.0125
Hess16HMC	0.1468	0.1468	-1.6e-05	0.4000	0.9994	3.8e-05
Hess32HMC	0.1461	0.1461	3.6e-05	0.8000	0.9993	3.7e-05
KFACLap	0.1456	0.1456	-2.2e-05	0.6000	0.9334	0.0002
LowRank32Lap	0.1358	0.1457	-0.0098	0.0000	0.7402	0.0098
LowRankLap	0.1351	0.1456	-0.0105	0.0000	0.7359	0.0105
RandHMC	0.1440	0.1440	-2.2e-05	0.2000	0.9766	6.8e-05
SGHMC move	0.1419	0.1457	-0.0037	0.0000	0.6796	0.0037
SGLD move	0.1425	0.1441	-0.0017	0.0000	0.7362	0.0017
SGLD-3	0.1368	0.1368	-4.9e-06	0.4000	0.9969	3.4e-05
SWAG	0.1361	0.1361	-5.4e-05	0.6000	0.9265	0.0003
SWAG20	0.1454	0.1455	-7.0e-05	0.2000	0.9528	0.0001
TrajHMC	0.2290	0.2292	-0.0002	0.4000	0.9915	0.0002
cSGLD move	0.1422	0.1454	-0.0033	0.0000	0.7046	0.0033
Fashion	0.2114	0.2122	-0.0008	0.0000	0.9212	0.0008
MNIST	0.2373	0.2385	-0.0011	0.0000	0.9298	0.0011

Table 7: Mode/ticket distribution-equivalence audit for selected posterior families. Delta is posterior-to-IMP minus matched chain-start-to-IMP. Wins is the fraction of seed-level paired posterior-chain deltas above zero; W-dist. is the two-sample Wasserstein distance over grouped seed means. A positive posterior mode account would require positive deltas rather than practical ties to chain-start magnitude supports.

Setting	Comparison	Modes	H_n	n_L	n_T	KS p	Ham.	Logit	Act.	H-cost	A-cost	Pass
Digits MLP	posterior samples vs tickets	1	0.0000	50	5	0.0788	0.6314	0.9819	n/a	0.0181	n/a	2/4
Digits MLP	posterior modes vs tickets	1	0.0000	1	5	0.9216	n/a	0.9821	n/a	0.0179	n/a	3/4
CIFAR subset ResNet	posterior samples vs tickets	1	0.0000	25	5	1.1e-05	0.6000	0.8878	n/a	0.1122	n/a	2/4
CIFAR subset ResNet	posterior modes vs tickets	1	0.0000	1	5	0.1388	n/a	0.9077	n/a	0.0923	n/a	3/4
CIFAR subset act. pilot	posterior samples vs tickets	3	0.4057	15	3	0.0330	0.6571	0.9022	0.8726	0.0978	0.1274	4/6
CIFAR subset act. pilot	posterior modes vs tickets	3	0.4057	3	3	0.2264	1.0000	0.9014	0.8704	0.0986	0.1296	6/6
CIFAR full ResNet	posterior samples vs tickets	1	0.0000	50	5	5.3e-09	0.0033	0.9369	0.9172	0.0631	0.0828	4/6
CIFAR full ResNet	posterior modes vs tickets	1	0.0000	1	5	0.0329	n/a	0.9383	0.9179	0.0617	0.0821	4/6
CIFAR full aligned	aligned posterior samples vs tickets	1	0.0000	50	5	2.3e-09	0.0000	0.9373	0.9168	0.0627	0.0832	4/6
CIFAR full aligned	aligned posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9363	0.9160	0.0637	0.0840	4/6
CIFAR full weight-aligned	weight-aligned chain start magnitude vs tickets	1	0.0000	5	5	4.8e-05	0.0000	0.9374	0.9195	0.0626	0.0805	4/6
CIFAR full weight-aligned	weight-aligned posterior samples vs tickets	1	0.0000	50	5	1.2e-08	0.1290	0.9336	0.9131	0.0664	0.0869	4/6
CIFAR full weight-aligned	weight-aligned posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9327	0.9132	0.0673	0.0868	4/6
CIFAR full cSGLD multi-chain	chain start magnitude vs tickets	1	0.0000	15	5	5.2e-07	0.0857	0.9363	0.9196	0.0637	0.0804	4/6
CIFAR full cSGLD multi-chain	posterior samples vs tickets	1	0.0000	75	5	3.3e-08	0.2461	0.9327	0.9144	0.0673	0.0856	4/6
CIFAR full cSGLD multi-chain	posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9288	0.9115	0.0712	0.0885	4/6
CIFAR full cSGLD independent	chain start magnitude vs tickets	1	0.0000	15	5	7.6e-08	0.0000	0.9306	0.9105	0.0694	0.0895	4/6
CIFAR full cSGLD independent	posterior samples vs tickets	1	0.0000	75	5	9.3e-10	0.0000	0.9269	0.9051	0.0731	0.0949	4/6
CIFAR full cSGLD independent	posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9242	0.9054	0.0758	0.0946	4/6
CIFAR full LowRank128Lap	chain start magnitude vs tickets	1	0.0000	5	5	9.0e-05	0.0000	0.9385	0.9218	0.0615	0.0782	4/6
CIFAR full LowRank128Lap	posterior samples vs tickets	1	0.0000	50	5	2.0e-06	0.8163	0.9319	0.9096	0.0681	0.0904	5/6
CIFAR full LowRank128Lap	posterior modes vs tickets	1	0.0000	1	5	0.1388	n/a	0.9289	0.9023	0.0711	0.0977	5/6
CIFAR full JointDiagLap270k	chain start magnitude vs tickets	1	0.0000	5	5	2.5e-05	0.0000	0.9389	0.9221	0.0611	0.0779	4/6
CIFAR full JointDiagLap270k	posterior samples vs tickets	1	0.0000	25	5	1.1e-08	0.0000	0.9373	0.9199	0.0627	0.0801	4/6
CIFAR full JointDiagLap270k	posterior modes vs tickets	1	0.0000	1	5	0.0413	n/a	0.9379	0.9225	0.0621	0.0775	4/6

Table 8: Direct proposal-level mode/ticket distribution probe on digits MLP plus CIFAR-10 ResNet-20 subset and full-data settings. Posterior sample masks and mean-shift posterior-mode representatives are compared with IMP tickets using the proposal’s layer-sparsity KS, mask-Hamming overlap, logit-space CKA, final-hidden-activation CKA, and Hungarian-cost thresholds; Modes and H_n report mean-shift cluster count and normalized raw-parameter basin entropy.

Source	Acc.	NLL	ECE	MSP AUROC	Ent. AUROC	MSP FPR95
Dense	0.8866	0.3536	0.0353	0.8230	0.8326	0.7636
IMP	0.8953	0.3387	0.0393	0.8306	0.8392	0.7508
SWAG ens.	0.8688	0.4018	0.0285	0.8050	0.8171	0.7715
Learned random	0.8449	0.4632	0.0270	0.7897	0.8024	0.7870
Gem-Miner	0.8418	0.4694	0.0283	0.7853	0.7985	0.7969
Var. prune	0.8301	0.5008	0.0255	0.7754	0.7897	0.8101

Table 9: Five-seed CIFAR-10 calibration and CIFAR-100 OOD diagnostics. SWAG ens. averages ten SWAG predictive samples per seed. Learned random, Gem-Miner, and Var. prune rows are trained from initialization at the IMP sparsity. MSP and entropy AUROC use maximum softmax probability and negative predictive entropy as in-distribution scores; lower FPR95 is better.

Source	Acc.	Support-IMP	Support-Dense	Support-Rewind
Epoch 1	0.4017	0.1782	0.2231	1.0000
Epoch 5	0.6881	0.2197	0.3582	0.3724
Epoch 10	0.7630	0.2342	0.5350	0.2746
Epoch 20	0.8492	0.2328	0.8433	0.2273
Epoch 30	0.8841	0.2312	1.0000	0.2231
max abs	–	0.2342	0.5616	0.3517
mean abs	–	0.2388	0.5314	0.4214
rms abs	–	0.2400	0.5713	0.4046

Table 10: Five-seed CIFAR dense-trajectory controls. The epoch-1 state is the IMP rewind point. Checkpoint rows use magnitude supports at that epoch; aggregate rows use trajectory score masks across checkpoints.

Source	Kind	Accuracy	Acc.-IMP	Support-IMP
IMP	IMP	0.8983	0.0006	1.0000
random 0	random	0.8422	-0.0555	0.0922
Epoch 1	checkpoint	0.8540	-0.0437	0.1785
Epoch 10	checkpoint	0.8730	-0.0247	0.2374
Epoch 30	checkpoint	0.8826	-0.0151	0.2356
rms abs	aggregate	0.8739	-0.0238	0.2429
mean abs	aggregate	0.8743	-0.0234	0.2421
path length	aggregate	0.8549	-0.0428	0.1772
rewind rms movement	aggregate	0.8572	-0.0405	0.1639
Gem-Miner	Gem-Miner	0.8471	-0.0499	0.0917
Var. prune	Var. prune	0.8306	-0.0669	0.0907
Hard concrete	Hard concrete	0.2766	-0.6204	0.0922

Table 11: Five-seed CIFAR trajectory-mask retraining controls. All masks are trained for 30 epochs from the same epoch-1 rewind state at the IMP sparsity. Acc.-IMP is candidate accuracy minus the final IMP accuracy from the corresponding seed.

Source	Kind	Acc.	ECE	Brier	Acc.-IMP	Supp.-IMP	Supp.-Dense
IMP	IMP	0.9711	0.0211	0.0410	0.0000	1.0000	0.7604
random 0	random	0.9489	0.0398	0.0821	-0.0222	0.2151	0.2139
Gem-Miner	Gem-Miner	0.9550	0.0392	0.0753	-0.0161	0.2081	0.2082
Var. prune	Var. prune	0.9633	0.0250	0.0532	-0.0078	0.3822	0.4010

Table 12: Five-seed digits MLP variational pruning sanity check. Var. prune optimizes Bernoulli mask probabilities on frozen initialization weights with Concrete samples, KL, sparsity, and entropy penalties before retraining the selected hard mask.

Base	Variant	Alpha	Accuracy	Acc.-IMP
Epoch 30	top IMP	0	0.8797	-0.0175
Epoch 30	top IMP	0.5	0.8882	-0.0090
Epoch 30	random non-IMP	0.5	0.8780	-0.0192
Epoch 30	top IMP	1	0.8966	-0.0006
rms abs	top IMP	0	0.8733	-0.0238
rms abs	top IMP	0.5	0.8851	-0.0121
rms abs	random non-IMP	0.5	0.8705	-0.0266
Epoch 10	top IMP	0	0.8712	-0.0259
Epoch 10	top IMP	0.5	0.8855	-0.0116
Epoch 10	random non-IMP	0.5	0.8704	-0.0268

Table 13: Five-seed CIFAR trajectory residual controls. Starting from a base trajectory mask, an IMP-residual row swaps the indicated fraction of base-only support for IMP-only support; a random-residual row swaps the same number of positions for non-IMP random support.

Base	Base	Low rm	Rand. rm	High rm	Rand. non
Epoch 30	0.8792	0.8881	0.8883	0.8906	0.8779
rms abs	0.8738	0.8874	0.8896	0.8914	0.8709
Epoch 10	0.8687	0.8862	0.8922	0.8920	0.8701

Table 14: Five-seed CIFAR residual removal-order controls. Low/Rand./High rm add the same top IMP-only residual support and vary which base-only weights are removed by low, random, or high base score. Rand. non removes low base-score weights but adds same-size non-IMP random support.

Base	Support-IMP	IMP-only(k)	Prune rnd.	Pred. AUC	Pred. recall
Epoch 30	0.2323	28.4k	2.9003	0.6197	0.2175
rms abs	0.2412	27.8k	2.9745	0.6165	0.2087
Epoch 10	0.2359	28.1k	2.9669	0.6206	0.2206

Table 15: Five-seed CIFAR residual anatomy. IMP-only(k) is the number of IMP-kept weights missing from the base support. Prune rnd. is the mean IMP pruning round for base-only weights. The predictor is a held-out logistic model over dense-trajectory rank features plus stage indicators, trained to recover IMP-only residual support among non-base weights.

Base	Base	Oracle	Pred.	Random	Pred. prec.	Rand. prec.
Epoch 30	0.8808	0.8892	0.8793	0.8805	0.1866	0.1253
rms abs	0.8740	0.8866	0.8744	0.8744	0.1834	0.1233
Epoch 10	0.8712	0.8890	0.8728	0.8724	0.1843	0.1239

Table 16: Five-seed CIFAR functional residual-predictor masks. For each base, the held-out predictor swaps half of base-only support for non-base weights with highest predicted IMP-only probability. Oracle uses true IMP-only residual support, and random uses same-size held-out non-IMP-random support. Prec. is the IMP-only precision of added weights.

Base	Base	Oracle	Cross-seed	Random	Cross prec.	Rand. prec.
Epoch 30	0.8815	0.8905	0.8776	0.8781	0.2413	0.1260
rms abs	0.8733	0.8878	0.8731	0.8725	0.2238	0.1246
Epoch 10	0.8730	0.8890	0.8745	0.8726	0.2388	0.1264

Table 17: Five-seed CIFAR cross-seed residual-transfer masks. For each held-out target seed, the predictor is trained on non-base candidate weights from the other four seeds, then swaps half of target base-only support for target non-base weights with highest predicted IMP-only probability. Oracle and random use the target seed only.

Base	Base	Oracle	Src vote	Aln src	Vote rand.	Aln rand.	Tgt rand.	Src prec.
Epoch 30	0.8813	0.8886	0.8769	0.8797	0.8774	0.8809	0.8779	0.1513
rms abs	0.8737	0.8872	0.8739	0.8727	0.8735	0.8743	0.8728	0.1451
Epoch 10	0.8720	0.8890	0.8725	0.8710	0.8732	0.8740	0.8719	0.1491

Table 18: Five-seed CIFAR direct cross-seed residual-support transfer. Source-vote variants add target non-base coordinates selected by votes from the other seeds’ oracle residual supports. Aln src and Aln rand. first map source channels to target channels by activation-correlation Hungarian matching. Tgt rand. is a target-seed random non-base residual control. Src prec. is target IMP-only precision among source-vote additions.

Base	T base	T oracle	T rand.	M base	M oracle	M rand.
Epoch 30	0.8827	0.8893	0.8791	0.8641	0.8926	0.8649
rms abs	0.8743	0.8892	0.8745	0.8605	0.8910	0.8628
Epoch 10	0.8744	0.8892	0.8712	0.8607	0.8942	0.8636

Table 19: Five-seed CIFAR residual base-compatibility controls. T rows use the actual trajectory base. M rows replace that base with a per-parameter random base preserving the same IMP/non-IMP counts and therefore the same base-to-IMP overlap before adding oracle or random residual support.

Base	M base	Top	Dense	RMS	RMS-d	Std	Rand	Dense ovlp	RMS ovlp	Rand ovlp
Epoch 30	0.8606	0.8915	0.8821	0.8852	0.8770	0.8710	0.8783	0.5573	0.5560	0.5006
rms abs	0.8631	0.8928	0.8812	0.8834	0.8745	0.8717	0.8795	0.5528	0.5510	0.5011
Epoch 10	0.8609	0.8911	0.8827	0.8829	0.8764	0.8710	0.8791	0.5565	0.5556	0.5002

Table 20: Five-seed CIFAR residual posterior-decomposition controls on IMP-overlap-matched random bases. Top adds the target oracle IMP-only residual. Dense, RMS, RMS-d, Std, and Rand add the same number of target IMP-only residual weights ranked by dense final magnitude, diagonal-Laplace posterior RMS, posterior RMS minus dense magnitude, posterior standard deviation, or uniform sampling. Ovlp columns report the fraction of additions that also belong to the top oracle additions.

Base	Base	Oracle	Rand. IMP	Global non	Param+score non
Epoch 30	0.8803	0.8872	0.8818	0.8767	0.8758
rms abs	0.8708	0.8858	0.8794	0.8701	0.8716
Epoch 10	0.8678	0.8854	0.8764	0.8695	0.8685

Table 21: Five-seed CIFAR residual stratified controls. All non-base variants remove the same low-base-score weights as the oracle. Rand. IMP adds random IMP-only residual weights. Global non adds random non-IMP non-base weights. Param+score non matches the oracle added weights by parameter tensor and within-parameter score decile, but excludes IMP-only weights.

Base	Base	Oracle	R1 acc.	R3 acc.	R5 acc.	R1 prec.	R3 prec.
Epoch 30	0.8793	0.8898	0.8792	0.8829	0.8867	0.4431	0.7523
rms abs	0.8736	0.8881	0.8771	0.8832	0.8857	0.4308	0.7459
Epoch 10	0.8724	0.8884	0.8772	0.8826	0.8861	0.4395	0.7536

Table 22: Five-seed CIFAR residual IMP-process probe. R1/R3/R5 rows add half residual support from IMP round-survivor candidates ranked by the corresponding round-trained weights. Prec. is the fraction of added weights that survive in the final IMP mask; R5 precision is one by construction.

Base	Rnd	Top acc.	Rand. acc.	Low acc.	Top final	Rand. final	Low final	Top oracle	Rand. oracle	Low oracle
Epoch 30	1	0.8806	0.8774	0.8776	0.4373	0.1913	0.0811	0.2706	0.0960	0.0340
Epoch 30	3	0.8839	0.8813	0.8770	0.7530	0.4357	0.1881	0.4777	0.2176	0.0652
Epoch 30	5	0.8887	0.8847	0.8815	1.0000	1.0000	1.0000	0.6818	0.4998	0.3182
rms abs	1	0.8788	0.8733	0.8738	0.4268	0.1871	0.0808	0.2617	0.0924	0.0343
rms abs	3	0.8830	0.8759	0.8726	0.7496	0.4294	0.1885	0.4734	0.2152	0.0657
rms abs	5	0.8857	0.8803	0.8759	1.0000	1.0000	1.0000	0.6803	0.4987	0.3197
Epoch 10	1	0.8795	0.8724	0.8699	0.4358	0.1880	0.0818	0.2698	0.0933	0.0345
Epoch 10	3	0.8822	0.8746	0.8727	0.7555	0.4347	0.1898	0.4795	0.2178	0.0657
Epoch 10	5	0.8855	0.8805	0.8755	1.0000	1.0000	1.0000	0.6821	0.4992	0.3179

Table 23: Five-seed CIFAR residual IMP-process ranking controls. Top uses the highest round-trained scores among round-survivor candidates; Rand. samples the same round-survivor candidate set; Low uses the lowest round-trained scores. Final is the fraction of added weights that survive in the final IMP mask; oracle is overlap with the final oracle added subset.

Base	Rnd	Score acc.	Match acc.	Δ	Wins	Oracle
Epoch 10	1	0.8823	0.8799	0.0024	5/5	0.5568
Epoch 10	3	0.8825	0.8789	0.0036	4/5	0.5976
Epoch 10	5	0.8851	0.8820	0.0030	5/5	0.6792
Epoch 30	1	0.8841	0.8843	-0.0002	3/5	0.5569
Epoch 30	3	0.8842	0.8831	0.0011	4/5	0.5971
Epoch 30	5	0.8860	0.8845	0.0014	3/5	0.6788
rms abs	1	0.8824	0.8787	0.0038	4/5	0.5538
rms abs	3	0.8821	0.8808	0.0013	3/5	0.5943
rms abs	5	0.8854	0.8836	0.0018	4/5	0.6773

Table 24: Five-seed CIFAR residual IMP-process oracle-overlap-matched control. Score acc. uses the round-trained score ordering inside the final-IMP residual candidate set. Match acc. samples random final-IMP residual additions with the same final-oracle overlap count. Δ is score minus matched random accuracy; wins are positive paired seed deltas.

Base	Rnd	Round acc.	Dense acc.	Base acc.	Δ_D	WinD	Δ_B	WinB	Oracle
Epoch 10	1	0.8833	0.8827	0.8824	0.0005	3/5	0.0008	4/5	0.5559
Epoch 10	3	0.8838	0.8827	0.8824	0.0011	4/5	0.0014	4/5	0.5984
Epoch 10	5	0.8872	0.8827	0.8824	0.0045	5/5	0.0048	5/5	0.6780
Epoch 30	1	0.8861	0.8831	0.8831	0.0030	5/5	0.0030	5/5	0.5574
Epoch 30	3	0.8874	0.8831	0.8831	0.0042	4/5	0.0042	4/5	0.5992
Epoch 30	5	0.8873	0.8831	0.8831	0.0041	5/5	0.0041	5/5	0.6779
rms abs	1	0.8820	0.8823	0.8820	-0.0004	2/5	-4.0e-05	2/5	0.5536
rms abs	3	0.8844	0.8823	0.8820	0.0021	4/5	0.0024	5/5	0.5950
rms abs	5	0.8870	0.8823	0.8820	0.0046	5/5	0.0050	5/5	0.6760

Table 25: Five-seed CIFAR residual IMP-process score-source controls. Candidate set is final-IMP residual for all variants; dense/base controls rank that candidate set by dense-final or base-source magnitudes. Δ_D/Δ_B are round-score minus dense/base score accuracy; wins are positive paired seed deltas.

Base	Rnd	Round acc.	Excl. acc.	Δ	Wins	Round oracle	Excl. oracle
Epoch 10	1	0.8844	0.8770	0.0074	5/5	0.5602	0.4398
Epoch 10	3	0.8858	0.8782	0.0076	5/5	0.5995	0.4005
Epoch 10	5	0.8872	0.8779	0.0093	5/5	0.6762	0.3238
Epoch 30	1	0.8848	0.8830	0.0018	4/5	0.5608	0.4392
Epoch 30	3	0.8867	0.8834	0.0033	5/5	0.5983	0.4017
Epoch 30	5	0.8867	0.8824	0.0043	5/5	0.6755	0.3245
rms abs	1	0.8846	0.8779	0.0067	5/5	0.5581	0.4419
rms abs	3	0.8858	0.8777	0.0081	5/5	0.5959	0.4041
rms abs	5	0.8861	0.8793	0.0068	5/5	0.6733	0.3267

Table 26: Five-seed CIFAR residual IMP-process round-exclusion control. Round acc. uses the round-trained score ordering inside the final-IMP residual candidate set. Excl. acc. removes the round-selected additions from that candidate set, then chooses the best remaining final-IMP residual additions by final IMP magnitude under the same support budget. Δ is round minus excluded accuracy; wins are positive paired seed deltas.

Base	Rnd	Round acc.	Excl. acc.	Tensor excl.	Δ_E	WinE	Δ_T	WinT	Tensor oracle
rms abs	5	0.8855	0.8753	0.8764	0.0103	5/5	0.0091	5/5	0.4153

Table 27: Five-seed CIFAR residual IMP-process tensor-matched round-exclusion control at the RMS trajectory base and round 5. Tensor excl. removes the round-selected additions, then chooses replacement final-IMP residual additions matched to the removed coordinates by parameter tensor before applying the final IMP magnitude score. Δ_E is round minus the unrestricted excluded oracle; Δ_T is round minus the tensor-matched excluded oracle.

Base	Rnd	Round acc.	Tensor excl.	Tensor+score excl.	Δ_T	WinT	Δ_{TS}	WinTS
rms abs	5	0.8878	0.8784	0.8837	0.0095	5/5	0.0041	5/5

Table 28: Five-seed CIFAR residual IMP-process tensor+score-matched round-exclusion control at the RMS trajectory base and round 5. Tensor+score excl. removes the round-selected additions, then chooses replacement final-IMP residual additions matched to the removed coordinates by parameter tensor and within-tensor round-score decile before applying final IMP magnitude. Δ_T is round minus the tensor-matched excluded oracle; Δ_{TS} is round minus the tensor+score-matched excluded oracle.

Base	Rnd	Round acc.	Resid. acc.	Dense acc.	Base acc.	Δ_R	WinR	Oracle drop
rms abs	5	0.8852	0.8811	0.8814	0.8811	0.0041	5/5	0.1830

Table 29: Five-seed CIFAR residual IMP-process projection control at the RMS trajectory base and round 5. Resid. acc. uses a round-score ordering after linearly projecting out the base-source, dense-final, and final-IMP magnitude scores inside the final-IMP residual candidate pool. Δ_R is round minus residualized-score accuracy; Oracle drop is round oracle overlap minus residualized oracle overlap.

Base	Rnd	Round acc.	Post-resid. acc.	Dense acc.	Base acc.	Δ_P	WinP	Oracle drop
rms abs	5	0.8847	0.8825	0.8815	0.8797	0.0023	5/5	0.1923

Table 30: Five-seed CIFAR residual IMP-process posterior-projection control at the RMS trajectory base and round 5. Post-resid. acc. uses a round-score ordering after linearly projecting out the base-source, dense-final, final-IMP magnitude, and diagonal-Laplace posterior RMS/std/excess scores inside the final-IMP residual candidate pool. Δ_P is round minus posterior-residualized accuracy; Oracle drop is round oracle overlap minus posterior-residualized oracle overlap.

Base	Rnd	Round acc.	Subspace-resid. acc.	Dense acc.	Base acc.	Δ_S	WinS	Oracle drop
rms abs	5	0.8869	0.8821	0.8822	0.8822	0.0048	5/5	0.1890

Table 31: Five-seed CIFAR residual IMP-process learned-subspace control at the RMS trajectory base and round 5. Subspace-resid. acc. uses a round-score ordering after projecting out the top rank-8 PCA component scores learned from dense trajectory, final-IMP magnitude, and earlier IMP-round scores inside the final-IMP residual candidate pool. Δ_S is round minus learned-subspace-residualized accuracy; Oracle drop is round oracle overlap minus learned-subspace-residualized oracle overlap.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes].

Justification: The abstract and Introduction state the support-equivalence claim, the negative conclusion, and the exact dense full-network CIFAR covariance limitation.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes].

Justification: The Limitations and Next Experiments section lists the remaining posterior, permutation, learned-mask, process-causality, and reproducibility limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete proof?

Answer: [N/A].

Justification: The paper is empirical and does not claim new theoretical results.

4. Experimental result reproducibility

Question: Does the paper fully disclose the information needed to reproduce the main experimental results?

Answer: [Yes].

Justification: The main text and generated appendix summarize the settings, and the accompanying artifact package contains scripts, run summaries, environment locks, a Docker artifact verifier, and a release manifest.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to reproduce the main experimental results?

Answer: [Yes].

Justification: The anonymized source repository is public at <https://github.com/evaldataset/lottery-ticket-bayesian-modes-artifact>, and the public release page hosts the artifact archive with SHA256 sidecar at <https://github.com/evaldataset/lottery-ticket-bayesian-modes-artifact/releases>. The local receipt registry records the exact archive URL, source commit, and passing public CI run; external GPU-container validation remains pending and is tracked separately as a release-readiness limitation.

6. Experimental setting/details

Question: Does the paper specify all training and test details necessary to understand the results?

Answer: [Yes].

Justification: The Operational Test and Current Results sections describe datasets, model families, posterior approximations, seeds, controls, and comparison metrics; the appendix and artifact manifest give generated tables and commands.

7. Experiment statistical significance

Question: Does the paper report error bars or other suitable information about statistical significance?

Answer: [Yes].

Justification: Main results are five-seed summaries with confidence intervals, paired win counts, grouped comparisons, and generated statistical tables.

523 **8. Experiments compute resources**

524 Question: For each experiment, does the paper provide sufficient information on compute
525 resources needed to reproduce the experiments?

526 Answer: [Yes].

527 Justification: The artifact package includes CPU and CUDA environment locks,
528 container definitions, exact reproduction command lines, dense-covariance mem-
529 ory bounds, artifact storage estimates, and compute-resource accounting in
530 docs/compute_resource_accounting.md. Historical metrics do not include profiler-
531 grade wall-clock fields, and that limitation is documented there.

532 **9. Code of ethics**

533 Question: Does the research conform with the NeurIPS Code of Ethics?

534 Answer: [Yes].

535 Justification: The work uses standard public benchmark datasets and synthetic/fake-data
536 smoke tests, does not involve human subjects, and reports negative results and limitations
537 transparently.

538 **10. Broader impacts**

539 Question: Does the paper discuss both potential positive and negative societal impacts of
540 the work?

541 Answer: [N/A].

542 Justification: This is foundational empirical analysis of pruning and posterior-support ex-
543 planations, with no direct deployed system or new capability; the main impact is method-
544 ological transparency around negative results.

545 **11. Safeguards**

546 Question: Does the paper describe safeguards for responsible release of data or models that
547 have a high risk for misuse?

548 Answer: [N/A].

549 Justification: The work does not release high-risk models, scraped datasets, generative
550 systems, or systems intended for deployment.

551 **12. Licenses for existing assets**

552 Question: Are existing assets properly credited and are license and terms of use explicitly
553 mentioned and respected?

554 Answer: [Yes].

555 Justification: Dataset, NeurIPS-style, and software dependency sources and license/terms
556 notes are inventoried in docs/asset_license_inventory.md. Raw benchmark datasets
557 are not redistributed in the local release manifest.

558 **13. New assets**

559 Question: Are new assets introduced in the paper well documented and provided alongside
560 the assets?

561 Answer: [Yes].

562 Justification: Project-created source, generated tables/figures, run summaries, saved
563 mask artifacts, reproducibility docs, container definitions, and excluded local
564 caches are documented in docs/new_asset_inventory.md and hash-inventoried by
565 docs/public_release_manifest.md.

566 **14. Crowdsourcing and research with human subjects**

567 Question: For crowdsourcing experiments and research with human subjects, does the pa-
568 per include the required details?

569 Answer: [N/A].

570 Justification: The paper does not use crowdsourcing or human-subject experiments.

571 **15. Institutional review board approvals**

572 Question: Does the paper describe IRB approvals or equivalent review for research with
573 human subjects?
574 Answer: [N/A].
575 Justification: The paper does not involve human-subject research.
576 **16. Declaration of LLM usage**
577 Question: Does the paper describe LLM usage if it is an important, original, or non-
578 standard component of the core methods?
579 Answer: [N/A].
580 Justification: LLMs are not part of the scientific method, experiments, posterior approxi-
581 mations, pruning algorithms, or analysis pipeline studied in the paper.